

# Kubernetes Advanced

## Agenda

### 1. Kubernetes - Überblick

- [Aufbau Allgemein](#)
- [Installation Variations](#)
- [Structure Kubernetes Deep Dive](#)
- [Ports und Protokolle](#)
- [kubelet garbage collection](#)

### 2. Kubernetes Controlplane

- [Renew Certificate](#)
- [HA-Cluster](#)

### 3. Installation

- [Basic Installation with microk8s](#)
- [Create cluster after basic installation](#)
- [Connect from remote](#)
- [Setup bash completion](#)

### 4. Kubernetes Imperative Commands (for debugging)

- [Example with run](#)

### 5. LoadBalancer

- [Setup internal loadBalancer for type LoadBalancer in Service](#)

### 6. Scale

- [Horizontal Pod Autoscaler](#)

### 7. Installation IngressController

- [Install Ingress On Digitalocean](#)

### 8. Kubernetes Praxis API-Objekte

- [Das Tool kubectl \(Devs/Ops\) - Cheatsheet](#)
- [Build applikation with Resource Objects](#)
- [Create Pod](#)
- [Create Replicaset](#)
- [kubectl/manifest/deployments](#)
- [Services - Aufbau](#)
- [kubectl/manifest/service](#)
- [DNS - Resolution - Services](#)
- [DaemonSets \(Devs/Ops\)](#)
- [Hintergrund Ingress](#)
- [Example with Hostnames](#)
- [Configmap MariaDB - Example](#)

### 9. Kubernetes - Sidecar Example

- [Kubernetes Sidecar](#)

### 10. Kubernetes - Probes

- [Überblick Probes](#)

### 11. Kubernetes - Wartung / Debugging

- [Create Network-Connection to pod](#)

### 12. Kubernetes Backup

- [Backups mit Velero](#)

### 13. Kubernetes Upgrade

- [Upgrade von tanzu \(Cluster API\)](#)

### 14. Monitoring with Prometheus / Grafana

- [Overview](#)
- [Setup prometheus/Grafana with helm](#)
- [exporters mongodb](#)
- [Good Kubernetes Board for Grafana](#)

### 15. Kubernetes Tipps & Tricks

- [kubectl kubeconfig mergen](#)
- [Create configmap from file](#)

### 16. Kubernetes Certificates (Control Plane) / Security

- [vmware - cluster api](#)
- [Pod Security Admission \(PSA\)](#)
- [Pod Security Policy \(PSP\)](#)

#### 17. Kubernetes Network / Firewall

- [Calico/Cilium - nginx example NetworkPolicy](#)
- [Egress / Ingress Examples with Exercise](#)
- [Mesh / istio](#)

#### 18. Kubernetes Probes (Liveness and Readiness)

- [Übung Liveness-Probe](#)
- [Übung Liveness http aus nginx](#)
- [Funktionsweise Readiness-Probe vs. Liveness-Probe](#)
- [Manueller Check readyz endpoint kubernetes api server aus pod](#)

#### 19. Kubernetes QoS / Limits / Requests

- [Quality of Service - evict pods](#)
- [Tools to identify LimitRange and Requests](#)

#### 20. Kubernetes Autoscaling

- [Autoscaling Pods/Deployments](#)

#### 21. Kubernetes Deployment Scenarios

- [Deployment green/blue, canary, rolling update](#)
- [Service Blue/Green](#)
- [Praxis-Übung A/B Deployment](#)

#### 22. Kubernetes Istio

- [Istio vs. Ingress Überblick](#)
- [Istio installieren und Addons bereitstellen](#)
- [Istion Überblick - egress und ingress - gateway](#)
- [Istio - Deployment of simple application](#)
- [Istio - Grafana Dashboard](#)

## Backlog

#### 1. Installation

- [Kubernetes mit der Cluster API aufsetzen](#)
- [Kubernetes mit kubadm aufsetzen \(calico\)](#)

#### 2. Kubernetes - Misc

- [Wann wird podIP vergeben ?](#)
- [Bash completion installieren](#)
- [Remote-Verbindung zu Kubernetes \(microk8s\) einrichten](#)
- [vim support for yaml](#)

#### 3. Kubernetes - Netzwerk (CNI's) / Mesh

- [Netzwerk Interna](#)
- [Übersicht Netzwerke](#)
- [IPV4/IPV6 Dualstack](#)
- [Ingress controller in microk8s aktivieren](#)

#### 4. Kubernetes - Ingress

- [ingress mit ssl absichern](#)

#### 5. Kubernetes - Wartung / Debugging

- [kubectl drain/uncordon](#)
- [Alte manifeste konvertieren mit convert plugin](#)
- [Curl from pod api-server](#)

#### 6. Kubernetes Praxis API-Objekte

- [kubectl example with run](#)
- [Ingress Controller auf Digitalocean \(doks\) mit helm installieren](#)
- [Documentation for default ingress nginx](#)
- [Beispiel Ingress](#)
- [Achtung: Ingress mit Helm - annotations](#)
- [Permanente Weiterleitung mit Ingress](#)
- [ConfigMap Example](#)
- [Configmap MariaDB my.cnf](#)

#### 7. Helm (Kubernetes Paketmanager)

- [Helm Grundlagen](#)
- [Helm Warum ?](#)
- [Helm Example](#)

## 8. Kubernetes - RBAC

- [Nutzer einrichten microk8s ab kubernetes 1.25](#)
- [Tipps&Tricks zu Deployment - Rollout](#)

## 9. Kustomize

- [Kustomize Overlay Beispiel](#)
- [Helm mit kustomize verheiraten](#)

## 10. Kubernetes - Tipps & Tricks

- [Kubernetes Debuggen ClusterIP/PodIP](#)
- [Debugging pods](#)
- [Taints und Tolerations](#)
- [pod aus deployment bei config - Änderung neu ausrollen](#)

## 11. Kubernetes Advanced

- [Curl api-server kubernetes aus pod heraus](#)

## 12. Kubernetes - Documentation

- [Documentation zu microk8s plugins/addons](#)
- [Shared Volumes - Welche gibt es ?](#)

## 13. Kubernetes - Hardening

- [Kubernetes Tipps Hardening](#)
- [Kubernetes Security Admission Controller Example](#)
- [Was muss ich bei der Netzwerk-Sicherheit beachten ?](#)

## 14. Kubernetes Interns / Misc.

- [OCI Container Images Standards](#)
- [Geolocation Kubernetes Cluster](#)

## 15. Kubernetes - Überblick

- [Installation - Welche Komponenten from scratch](#)

## 16. Kubernetes - microk8s (Installation und Management)

- [kubect! unter windows - Remote-Verbindung zu Kuberens \(microk8s\) einrichten](#)
- [Arbeiten mit der Registry](#)
- [Installation Kubernetes Dashboard](#)

## 17. Kubernetes - RBAC

- [Nutzer einrichten - kubernetes bis 1.24](#)

## 18. kubect!

- [Tipps&Tricks zu Deployment - Rollout](#)

## 19. Kubernetes - Monitoring (microk8s und vanilla)

- [metrics-server aktivieren \(microk8s und vanilla\)](#)

## 20. Kubernetes - Backups

- [Kubernetes Aware Cloud Backup - kasten.io](#)

## 21. Kubernetes - Tipps & Tricks

- [Assigning Pods to Nodes](#)

## 22. Kubernetes - Documentation

- [LDAP-Anbindung](#)
- [Helpful to learn - Kubernetes](#)
- [Environment to learn](#)
- [Environment to learn II](#)
- [Youtube Channel](#)

## 23. Kubernetes - Shared Volumes

- [Shared Volumes with nfs](#)

## 24. Kubernetes - Hardening

- [Kubernetes Tipps Hardening](#)

# Kubernetes - Überblick

## Aufbau Allgemein

### Architecture



## Komponenten / Grundbegriffe

### Master (Control Plane)

#### Aufgaben

- Der Master koordiniert den Cluster
- Der Master koordiniert alle Aktivitäten in Ihrem Cluster
  - Planen von Anwendungen
  - Verwalten des gewünschten Status der Anwendungen
  - Skalieren von Anwendungen
  - Rollout neuer Updates.

### Komponenten des Masters

#### ETCD

- Verwalten der Konfiguration des Clusters (key/value - pairs)

#### KUBE-CONTROLLER-MANAGER

- Zuständig für die Überwachung der Stati im Cluster mit Hilfe von endlos loops.
- kommuniziert mit dem Cluster über die kubernetes-api (bereitgestellt vom kube-api-server)

#### KUBE-API-SERVER

- provides api-frontent for administration (no gui)
- Exposes an HTTP API (users, parts of the cluster and external components communicate with it)
- REST API

#### KUBE-SCHEDULER

- assigns Pods to Nodes.
- scheduler determines which Nodes are valid placements for each Pod in the scheduling queue ( according to constraints and available resources )
- The scheduler then ranks each valid Node and binds the Pod to a suitable Node.
- Reference implementation (other schedulers can be used)

### Nodes

- Nodes (Knoten) sind die Arbeiter (Maschinen), die Anwendungen ausführen
- Ref: <https://kubernetes.io/de/docs/concepts/architecture/nodes/>

### Pod/Pods

- Pods sind die kleinsten einsetzbaren Einheiten, die in Kubernetes erstellt und verwaltet werden können.
- Ein Pod (übersetzt Gruppe) ist eine Gruppe von einem oder mehreren Containern
  - gemeinsam genutzter Speicher- und Netzwerkressourcen
  - Befinden sich immer auf dem gleich virtuellen Server

## Control Plane Node (former: master) - components

## Node (Minion) - components

### General

- On the nodes we will rollout the applications

### kubelet

```
Node Agent that runs on every node (worker)
Er stellt sicher, dass Container in einem Pod ausgeführt werden.
```

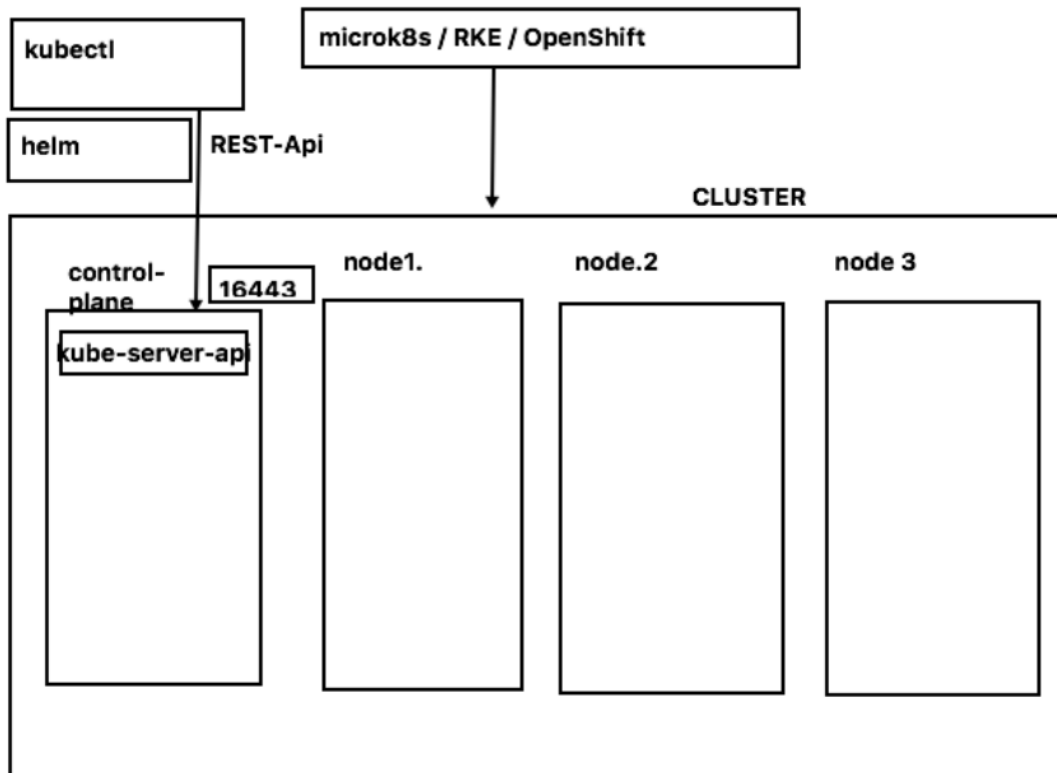
### Kube-proxy

- Läuft auf jedem Node
- = Netzwerk-Proxy für die Kubernetes-Netzwerk-Services.
- Kube-proxy verwaltet die Netzwerkkommunikation innerhalb oder außerhalb Ihres Clusters.

### Referenzen

- <https://www.redhat.com/en/topics/containers/kubernetes-architecture>

## Installation Variations



## Structure Kubernetes Deep Dive

- <https://github.com/jmetzger/training-kubernetes-advanced/assets/1933318/1ca0d174-f354-43b2-81cc-67af8498b56c>

## Ports und Protokolle

- <https://kubernetes.io/docs/reference/networking/ports-and-protocols/>

## kubelet garbage collection

### What is do ?

- Deletes unused containers after 1 minutes
- and unused images after 5 minutes

### Reference:

- <https://kubernetes.io/docs/concepts/architecture/garbage-collection/#containers-images>

## list images with ctr

- ctr is the cli tool for containerd

```
## from client
kubectl run nginx --image nginx

## on worker - node
ctr images list | grep nginx
```

## Kubernetes Controlplane

### Renew Certificate

### Zertifikate überprüfen

```
kubeadm certs check-expiration
```

. Wo werden Zertifikate benötigt ?

- zum kube-apiserver hin von den einzelnen Komponenten

```
- zum  
usw.
```

## Sonderrolle

b. Sonderrolle kubelet

Macht ein automatisches Renew the certificate über die  
Zertifikat api. Schritte:

Es erfolgt ein automatisches Approval des Signing Requests  
über den Controller Manager

Diese muss aktiviert sein:

```
https://kubernetes.io/docs/tasks/tls/certificate-rotation/  
--rotate-certificates
```

```
root@worker1:/var/lib/kubelet# grep -r "rotate" config.yaml  
rotateCertificates: true
```

## Zertifikatserneuerung

### Schritt 1:

c. Wir erneuern wir Zertifikate ?

Wichtig: Das muss auf allen Control-Nodes passieren, wenn sie kurz vor dem ablaufen sind.

auf dem controlplane (bspw. api-server)  
kubeadm certs renew apiserver

### Schritt 2:

```
## nochmal gucken, welches Zertifikat genommen  
echo | openssl s_client -showcerts -connect 64.226.76.200:6443 -servername api 2>/dev/null | openssl x509 -noout -enddate  
  
### Wichtig, kein kubectl delete po verwenden .  
## command output may be misleading in describing static pods: even if it shows that the static pod restarted recently, the  
correspondent pod containers were not restarted.  
  
## dann das manifests wegschieben  
cd /etc/kubernetes/manifests/  
mv kube-apiserver.yaml /tmp  
  
## will not work anymore, because apiserver is not running  
kubectl -n kube-system get pods
```

### Schritt 3: mit low-level tools checken pod noch läuft / wieder läuft

```
export CONTAINER_RUNTIME_ENDPOINT=unix:///var/run/containerd/containerd.sock  
## taucht nicht mehr auf -> apiservre  
crictl pods
```

```
mv /tmp/kube-apiserver.yaml .  
crictl pods | grep api
```

```
kubectl get nodes
```

```
kubeadm certs check-expiration | grep "apiserver "  
apiserver          Jan 23, 2025 04:09 UTC    364d          ca          no
```

```
echo | openssl s_client -showcerts -connect 64.226.76.200:6443 -servername api 2>/dev/null | openssl x509 -noout -enddate  
notAfter=Jan 23 04:09:04 2025 GMT
```

## Zertifikate ohne Downtime

Das wird nur funktionieren, wenn mir eine HA-Cluster haben. Dort gibt es mehrere Controlplanes und wir haben einen LoadBalancer davor. -> hier vielleicht noch ein Schaubild zeigen.

Ansonsten muss immer der kube-api-server neu gestartet werden und die einzelnen Komponenten, hier haben wir immer eine kurze Downtime.

Dies wird durch ein HA-Cluster vermieden. Dort ist ein LoadBalancer davorgeschaltet.

```
### HA-Cluster

### Übersicht

! [image] (https://github.com/jmetzger/training-kubernetes-advanced/assets/1933318/9f791d15-8c97-4f07-862b-cc2bf6035dc0)

### Aufsetzen eines HA-Clusters (auf vm's oder Metall)

* https://kubernetes.io/docs/v3.4/installing-on-linux/high-availability-configurations/set-up-ha-cluster-using-keepalived-haproxy/
* https://mvallim.github.io/kubernetes-under-the-hood/documentation/haproxy-cluster.html
* https://www.lisenet.com/2021/install-and-configure-a-multi-master-ha-kubernetes-cluster-with-kubeadm-haproxy-and-keepalived-on-centos-7/

### Aufsetzen eines HA-Cluster (Internal)

* https://github.com/kubernetes/kubekey/blob/master/docs/ha-mode.md

### Varianten den LoadBalancer zu platzieren

* https://github.com/kubernetes/kubeadm/blob/main/docs/ha-considerations.md

## Installation

### Basic Installation with microk8s

### Walkthrough
```

```
sudo snap install microk8s --classic microk8s status
```

```
### Optional
```

## Execute kubectl commands like so

```
microk8s kubectl microk8s kubectl cluster-info
```

## Make it easier with an alias

```
echo "alias kubectl='microk8s kubectl'" >> ~/.bashrc source ~/.bashrc kubectl
```

```
### Working with snaps
```

```
snap info microk8s
```

```
### Ref:

* https://microk8s.io/docs/setting-snap-channel

### Create cluster after basic installation

### Walkthrough
```

## auf master (jeweils für jedes node neu ausführen)

```
microk8s add-node
```

## dann auf jeweiligem node vorigen Befehl der ausgegeben wurde ausführen

## Kann mehr als 60 sekunden dauern ! Geduld...Geduld..Geduld

```
##z.B. -> ACHTUNG evtl. IP ändern microk8s join 10.128.63.86:25000/567a21bdfc9a64738ef4b3286b2b8a69
```

```
### Auf einem Node addon aktivieren z.B. ingress
```

gucken, ob es auf dem anderen node auch aktiv ist.

```
### Ref:
```

```
* https://microk8s.io/docs/high-availability

### Connect from remote

### Step 1: Install kubectl on local machine (or jump-server)
```

## on CLIENT install kubectl

```
sudo snap install kubectl --classic
```

```
### Step 2: configure kubectl
```

## On MASTER -server get config

### als root

```
microk8s config > /tmp/config cat /tmp/config
```

## Optional or simply copy & paste

## Download (scp config file) and store in .kube - folder

```
cd mkdir .kube cd .kube # Wichtig: config muss nachher im verzeichnis .kube liegen
```

## scp kurs@master\_server:/path/to/remote\_config config

### z.B.

```
scp kurs@192.168.56.102:/home/kurs/remote_config config
```

## oder benutzer 11trainingdo

```
scp 11trainingdo@192.168.56.102:/home/11trainingdo/remote_config config
```

Evtl. IP-Adresse in config zum Server aendern

## Ultimative 1. Test auf CLIENT

```
kubectl cluster-info
```

## or if using kubectl or alias

```
kubectl get pods
```

## if you want to use a different kube config file, you can do like so

```
kubectl --kubeconfig /home/myuser/.kube/myconfig
```

```
### Setup bash completion
```

```
### Walkthrough
```

```
apt install bash-completion source /usr/share/bash-completion/bash_completion
```

## is it installed properly

```
type _init_completion
```

## activate for all users

```
kubectl completion bash | sudo tee /etc/bash_completion.d/kubectl > /dev/null
```

## verify with new login

## example



su - tln1

## zum Testen

kubect! g

```
## Kubernetes Imperative Commands (for debugging)

### Example with run

### Example (that does work)
```

## Show the pods that are running

kubect! get pods

## Synopsis (most simplistic example

**kubect! run NAME --image=IMAGE\_EG\_FROM\_DOCKER**

## example

kubect! run nginx --image=nginx:1.23

kubect! get pods

## on which node does it run ?

kubect! get pods -o wide

```
### Example (that does not work)
```

kubect! run foo2 --image=foo2

## ImageErrPull - Image konnte nicht geladen werden

kubect! get pods

## Weitere status - info

kubect! describe pods foo2

```
### Ref:

* https://kubernetes.io/docs/reference/generated/kubect! /kubect!-commands#run

## LoadBalancer

### Setup internal loadBalancer for type LoadBalancer in Service

### Step 1: helm chart install
```

helm repo add metallb <https://metallb.github.io/metallb> helm install metallb metallb/metallb --version 0.14.5 --namespace=metallb-system --create-namespace

## Wait for all the systems to come up

kubect! -n metallb-system get pods -o wide

```
### Step 2: addresspool und Propagation-type (config)
```

cd mkdir -p manifests cd manifests mkdir lb cd lb vi 01-addresspool.yml

apiVersion: metallb.io/v1beta1 kind: IPAddressPool metadata: name: first-pool namespace: metallb-system spec: addresses:

## we will use our external ip here

- 134.209.231.154-134.209.231.154

## both notations are possible

- 157.230.113.124/32

```
kubectl apply -f .
```

```
vi 02-advertisement.yml
```

```
apiVersion: metallb.io/v1beta1 kind: L2Advertisement metadata: name: example namespace: metallb-system
```

```
kubectl apply -f .
```

```
### Schritt 4: Test do i get an external ip
```

```
vi 03-deploy.yml
```

```
apiVersion: apps/v1 kind: Deployment metadata: name: my-nginx spec: selector: matchLabels: run: web-nginx replicas: 3 template: metadata: labels: run: web-nginx spec: containers: - name: cont-nginx image: nginx ports: - containerPort: 80
```

```
vi 04-service.yml
```

### 02-svc.yml

```
apiVersion: v1 kind: Service metadata: name: svc-nginx labels: svc: nginx spec: type: LoadBalancer ports:
```

- port: 80 protocol: TCP selector: run: web-nginx

```
kubectl apply -f . kubectl get pods kubectl get svc
```

```
kubectl delete -f 03-deploy.yml 04-service.yml
```

## Scale

### Horizontal Pod Autoscaler

**Example: newest version with autoscaling/v2 used to be hpa/v1**

#### Prerequisites

- Metrics-Server needs to be running

```
## Test with
kubectl top pods
```

```
## Install
kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
## after that it will be available in kube-system namespace as pod
kubectl -n kube-system get pods | grep -i metrics
```

#### Step 1: deploy app

```
cd
mkdir -p manifests
cd manifests
mkdir hpa
cd hpa
vi 01-deploy.yaml
```

```

---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello
spec:
  replicas: 3
  selector:
    matchLabels:
      app: hello
  template:
    metadata:
      labels:
        app: hello
    spec:
      containers:
        - name: hello
          image: k8s.gcr.io/hpa-example
          resources:
            requests:
              cpu: 100m
---
kind: Service
apiVersion: v1
metadata:
  name: hello
spec:
  selector:
    app: hello
  ports:
    - port: 80
      targetPort: 80
---
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: hello
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: hello
  minReplicas: 2
  maxReplicas: 20
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 80

```

## Step 2: Load Generator

```
vi 02-loadgenerator.yml
```

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: load-generator
  labels:
    app: load-generator
spec:
  replicas: 100
  selector:
    matchLabels:
      app: load-generator
  template:
    metadata:
      name: load-generator
      labels:
        app: load-generator
    spec:
      containers:
        - name: load-generator
          image: busybox
          command:

```

```
- /bin/sh
- -c
- "while true; do wget -q -O- http://hello.default.svc.cluster.local; done"
```

## Downscaling

- Downscaling will happen after 5 minutes o

```
## Adjust down to 1 minute
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: hello
spec:
  # change to 60 secs here
  behavior:
    scaleDown:
      stabilizationWindowSeconds: 60
  # end of behaviour change
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: hello
  minReplicas: 2
  maxReplicas: 20
  metrics:
  - type: Resource
    resource:
      name: cpu
      target:
        type: Utilization
        averageUtilization: 80
```

For scaling down the stabilization window is 300 seconds (or the value of the `--horizontal-pod-autoscaler-downscale-stabilization` flag if provided)

## Prevent Downscaling

```
## Adjust down to 1 minute
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: hello
spec:
  # change to 60 secs here
  behavior:
    scaleDown:
      selectPolicy: Disabled
  # end of behaviour change
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: hello
  minReplicas: 2
  maxReplicas: 20
  metrics:
  - type: Resource
    resource:
      name: cpu
      target:
        type: Utilization
        averageUtilization: 80
```

## Reference

- <https://docs.digitalocean.com/tutorials/cluster-autoscaling-ca-hpa/>
- <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale-walkthrough/#autoscaling-on-more-specific-metrics>
- <https://medium.com/expedia-group-tech/autoscaling-in-kubernetes-why-doesnt-the-horizontal-pod-autoscaler-work-for-me-5f0094694054>

## Installation IngressController

### Install Ingress On Digitalocean

#### Basics

- works for all platform with helm if no ingresscontroller ist present

- if you have ingress - objects and no ingresscontroller nothing works

## Prerequisites

- kubectl must be set up

## Walkthrough (Setup Ingress Controller)

```
helm repo add ingress-nginx https://kubernetes.github.io/ingress-nginx
helm repo update

## helm show values ingress-nginx/ingress-nginx

## It will be setup with type loadbalancer - so waiting to retrieve an ip from the external loadbalancer
## This will take a little.
helm install nginx-ingress ingress-nginx/ingress-nginx --namespace ingress --create-namespace --set
controller.publishService.enabled=true

## See when the external ip comes available
kubectl -n ingress get all
kubectl --namespace ingress get services -o wide -w nginx-ingress-ingress-nginx-controller

## Output
NAME                                TYPE                CLUSTER-IP      EXTERNAL-IP      PORT(S)              AGE
SELECTOR
nginx-ingress-ingress-nginx-controller  LoadBalancer      10.245.78.34    157.245.20.222   80:31588/TCP,443:30704/TCP  4m39s
app.kubernetes.io/component=controller,app.kubernetes.io/instance=nginx-ingress,app.kubernetes.io/name=ingress-nginx

## Now setup wildcard - domain for training purpose
## inwx.com
*.lab1.t3isp.de A 157.245.20.222
```

## Kubernetes Praxis API-Objekte

### Das Tool kubectl (Devs/Ops) - Cheatsheet

#### Allgemein

```
## Zeige Information über das Cluster
## Show Information about the cluster
kubectl cluster-info

kubectl get nodes
kubectl get nodes -o wide

## Which api-resources ?
kubectl api-resources

## Hilfe zu object und eigenschaften bekommen
kubectl explain pod
kubectl explain pod.metadata
kubectl explain pod.metadata.name
```

#### Arbeiten mit manifesten

```
kubectl apply -f nginx-replicaset.yml
## Wie ist aktuell die hinterlegte config im system
kubectl get -o yaml -f nginx-replicaset.yml

## Änderung in nginx-replicaset.yml z.B. replicas: 4
## dry-run - was wird geändert
kubectl diff -f nginx-replicaset.yml

## anwenden
kubectl apply -f nginx-replicaset.yml

## Alle Objekte aus manifest löschen
kubectl delete -f nginx-replicaset.yml
```

#### Ausgabeformate

```
## Ausgabe kann in verschiedenen Formaten erfolgen
kubectl get pods -o wide # weitere informationen
## im json format
kubectl get pods -o json
```

```
## gilt natürluch auch für andere kommandos
kubectl get deploy -o json
kubectl get deploy -o yaml

## get a specific value from the complete json - tree
kubectl get node k8s-nue-jo-ff1p1 -o=jsonpath='{.metadata.labels}'
```

## Zu den Pods

```
## Start einen pod // BESSER: direkt manifest verwenden
## kubectl run podname image=imagename
kubectl run nginx image=nginx

## Pods anzeigen
kubectl get pods
kubectl get pod
## Format weitere Information
kubectl get pod -o wide
## Zeige labels der Pods
kubectl get pods --show-labels

## Zeige pods mit einem bestimmten label
kubectl get pods -l app=nginx

## Status eines Pods anzeigen
kubectl describe pod nginx

## Pod löschen
kubectl delete pod nginx

## Kommando in pod ausführen
kubectl exec -it nginx -- bash
```

## Arbeiten mit namespaces

```
## Welche namespaces auf dem System
kubectl get ns
kubectl get namespaces
## Standardmäßig wird immer der default namespace verwendet
## wenn man kommandos aufruft
kubectl get deployments

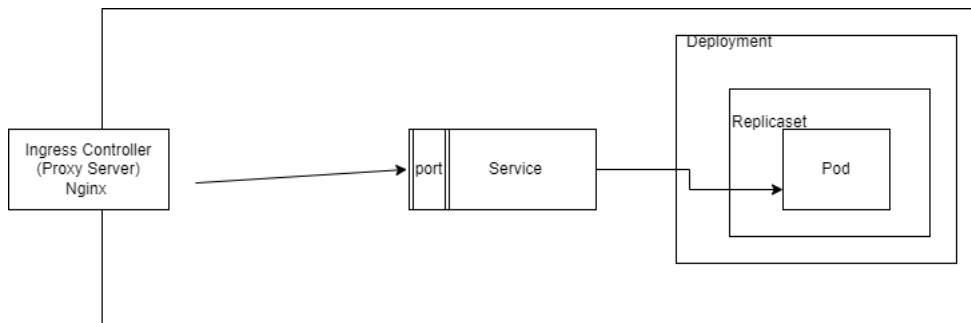
## Möchte ich z.B. deployment vom kube-system (installation) aufrufen,
## kann ich den namespace angeben
kubectl get deployments --namespace=kube-system
kubectl get deployments -n kube-system

## wir wollen unseren default namespace ändern
kubectl config set-context --current --namespace <dein-namespace>
```

## Referenz

- <https://kubernetes.io/de/docs/reference/kubectl/cheatsheet/>

## Build applikation with Resource Objects



## Create Pod

## Walkthrough

```
cd
mkdir -p manifests
cd manifests
mkdir -p web
cd web
```

```
## vi nginx-static.yml

apiVersion: v1
kind: Pod
metadata:
  name: nginx-static-web
  labels:
    webserver: nginx
spec:
  containers:
    - name: web
      image: nginx:1.23
```

```
kubectl apply -f nginx-static.yml
kubectl describe pod nginx-static-web
## show config
kubectl get pod/nginx-static-web -o yaml
kubectl get pod/nginx-static-web -o wide
```

## Create Replicaset

```
cd
mkdir -p manifests
cd manifests
mkdir 02-rs
cd 02-rs
vi rs.yml
```

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: nginx-replica-set
spec:
  replicas: 2
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      name: template-nginx-replica-set
      labels:
        tier: frontend
    spec:
      containers:
        - name: nginx
          image: nginx:1.21
          ports:
            - containerPort: 80
```

```
kubectl apply -f .
kubectl get all
## delete one of the pods
kubectl delete po nginx-replica-set-xyzhg
## new pod should have been created
kubectl get all
```


## kubectl/manifest/deployments

```
cd
mkdir -p manifests
cd manifests
mkdir 03-deploy
cd 03-deploy
nano deploy.yml
```

```
## vi deploy.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 8 # tells deployment to run 8 pods matching the template
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.21
          ports:
            - containerPort: 80
```

```
kubectl apply -f .
kubectl get all
```

## Services - Aufbau



## kubectl/manifest/service

### Schritt 1: Deployment

```
cd
mkdir -p manifests
cd manifests
mkdir 04-service
cd 04-service

## 01-deploy.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-nginx
spec:
  selector:
    matchLabels:
      run: my-nginx
  replicas: 3
  template:
    metadata:
      labels:
        run: my-nginx
    spec:
      containers:
        - name: my-nginx
          image: nginx
          ports:
            - containerPort: 80
```

```
kubectl apply -f .
```

### Schritt 2:



```
## 02-svc.yml
apiVersion: v1
kind: Service
metadata:
  name: my-nginx
spec:
  ports:
    - port: 80
      protocol: TCP
  selector:
    run: my-nginx
```

```
kubectl apply -f .
kubectl get svc my-nginx
kubectl describe svc my-nginx
```

```
## Testing
kubectl delete -f 01-deploy.yml

## No endpoints in svc
kubectl describe svc my-nginx
kubectl apply -f 01-deploy.yml
kubectl describe svc my-nginx
```

## Schritt 2b: NodePort

```
## 02-svc.yml
apiVersion: v1
kind: Service
metadata:
  name: my-nginx
  labels:
    svc: nginx
spec:
  type: NodePort
  ports:
    - port: 80
      protocol: TCP
  selector:
    run: my-nginx
```

```
kubectl apply -f .
```

## Ref.

- <https://kubernetes.io/docs/concepts/services-networking/connect-applications-service/>

## DNS - Resolution - Services

### How does it work

```
3 Variants:

svc-name
or:
svc-name.<namespace>
or:
svc-name.<namespace>.svc.cluster.local
```

## Example

```
kubectl run podtest --rm -ti --image busybox
If you don't see a command prompt, try pressing enter.
/ # wget -O - http://apple-service.jochen
Connecting to apple-service.jochen (10.245.39.214:80)
writing to stdout
apple-tln1
-
100%
|*****| 11 0:00:00
ETA
written to stdout
/ # wget -O - http://apple-service.jochen.svc.cluster.local
Connecting to apple-service.jochen.svc.cluster.local (10.245.39.214:80)
writing to stdout
apple-tln1
-
100%
```

```
|*****| 11 0:00:00
ETA
written to stdout
/ # wget -O - http://apple-service
Connecting to apple-service (10.245.39.214:80)
writing to stdout
apple-tln1
-
100%
|*****| 11 0:00:00
ETA
written to stdout
```

## Hintergrund Ingress

### Ref. / Dokumentation

- <https://matthewpalmer.net/kubernetes-app-developer/articles/kubernetes-ingress-guide-nginx-example.html>

### Example with Hostnames

#### Walkthrough

##### Step 1: pods and services

```
cd
mkdir -p manifests
cd manifests
mkdir abi
cd abi
```

```
## apple.yml
## vi apple.yml
kind: Pod
apiVersion: v1
metadata:
  name: apple-app
  labels:
    app: apple
spec:
  containers:
    - name: apple-app
      image: hashicorp/http-echo
      args:
        - "-text=apple-<dein-name>"
---

kind: Service
apiVersion: v1
metadata:
  name: apple-service
spec:
  selector:
    app: apple
  ports:
    - protocol: TCP
      port: 80
      targetPort: 5678 # Default port for image
```

```
kubect1 apply -f apple.yml
```

```
## banana
## vi banana.yml
kind: Pod
apiVersion: v1
metadata:
  name: banana-app
  labels:
    app: banana
spec:
  containers:
    - name: banana-app
      image: hashicorp/http-echo
      args:
        - "-text=banana-<dein-name>"
---

kind: Service
```

```
apiVersion: v1
metadata:
  name: banana-service
spec:
  selector:
    app: banana
  ports:
    - port: 80
      targetPort: 5678 # Default port for image
```

```
kubectl apply -f banana.yml
```

## Step 2: Ingress

```
## Ingress
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: example-ingress
  annotations:
    ingress.kubernetes.io/rewrite-target: /
    # with the ingress controller from helm, you need to set an annotation
    # otherwise it does not know, which controller to use
    # old version... use ingressClassName instead
    # kubernetes.io/ingress.class: nginx
spec:
  ingressClassName: nginx
  rules:
    - host: "<evername>.lab<nr>.t3isp.de"
      http:
        paths:
          - path: /apple
            backend:
              serviceName: apple-service
              servicePort: 80
          - path: /banana
            backend:
              serviceName: banana-service
              servicePort: 80
```

```
## ingress
kubectl apply -f ingress.yml
kubectl get ing
```

## Reference

- <https://matthewpalmer.net/kubernetes-app-developer/articles/kubernetes-ingress-guide-nginx-example.html>

## Find the problem

```
## Hints

## 1. Which resources does our version of kubectl support
## Can we find Ingress as "Kind" here.
kubectl api-resources

## 2. Let's see, how the configuration works
kubectl explain --api-version=networking.k8s.io/v1 ingress.spec.rules.http.paths.backend.service

## now we can adjust our config
```

## Solution

```
## in kubernetes 1.22.2 - ingress.yml needs to be modified like so.
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example-ingress
  annotations:
    ingress.kubernetes.io/rewrite-target: /
    # with the ingress controller from helm, you need to set an annotation
    # old version useClassName instead
    # otherwise it does not know, which controller to use
    # kubernetes.io/ingress.class: nginx
spec:
  ingressClassName: nginx
  rules:
```

```
- host: "app12.lab.t3isp.de"
http:
  paths:
    - path: /apple
      pathType: Prefix
      backend:
        service:
          name: apple-service
          port:
            number: 80
    - path: /banana
      pathType: Prefix
      backend:
        service:
          name: banana-service
          port:
            number: 80
```

## Configmap MariaDB - Example

### Step 1: configmap

```
cd
mkdir -p manifests
cd manifests
mkdir cftest
cd cftest
nano 01-configmap.yml
```

```
### 01-configmap.yml
kind: ConfigMap
apiVersion: v1
metadata:
  name: mariadb-configmap
data:
  # als Wertepaare
  MARIADB_ROOT_PASSWORD: 11abc432
```

```
kubectl apply -f .
kubectl get cm
kubectl get cm mariadb-configmap -o yaml
```

### Step 2: Deployment

```
nano 02-deploy.yml
```

```
##deploy.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mariadb-deployment
spec:
  selector:
    matchLabels:
      app: mariadb
  replicas: 1
  template:
    metadata:
      labels:
        app: mariadb
    spec:
      containers:
        - name: mariadb-cont
          image: mariadb:latest
          envFrom:
            - configMapRef:
                name: mariadb-configmap
```

```
kubectl apply -f .
kubectl exec -it mariadb-deployment-c6df6f959-9jvkb -- bash
```

```
## env
env
env | grep -i mariadb_root
```

### Important Sidenode

- If configmap changes, deployment does not know
- So kubectl apply -f deploy.yml will not have any effect
- to fix, use stakater/reloader: <https://github.com/stakater/Reloader>

## Kubernetes - Sidecar Example

### Kubernetes Sidecar

#### Walkthrough

```
cd
mkdir -p manifests
cd manifests
mkdir -p sidecar
cd sidecar
```

```
nano 01-pod.yml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: sidecar-example
spec:
  securityContext:
    runAsUser: 0
    runAsGroup: 0
  containers:
    - name: splunk-uf
      image: splunk/universalforwarder:latest
      env:
        - name: SPLUNK_START_ARGS
          value: --accept-license
        - name: SPLUNK_USER
          value: root
        - name: SPLUNK_GROUP
          value: root
        - name: SPLUNK_PASSWORD
          value: helloworld
        - name: SPLUNK_CMD
          value: add monitor /var/log/
        - name: SPLUNK_STANDALONE_URL
          value: splunk.company.internal
      volumeMounts:
        - name: shared-data
          mountPath: /var/log
    - name: my-nginx
      image: nginx
      volumeMounts:
        - name: shared-data
          mountPath: /var/log/nginx/
  volumes:
    - name: shared-data
      emptyDir: {}
```

```
kubectl apply -f .
kubectl get pods sidecar-example
kubectl describe pods sidecar-example
## exec into my-nginx
kubectl exec -it sidecar-example -c my-nginx -- sh
```

## Kubernetes - Probes

### Überblick Probes

#### Welche Probes gibt es ?

- startup (probe)
- liveness (probe)
- readiness (probe)

#### Wo werden die Probes definiert ?

- Die Probes werden immer auf Container-Ebene definiert

#### Liveness Probe

#### Was ist das Standardverhalten (wenn keine Liveness Probe existiert)

- Es muss ein Prozess mit der id 1 laufen (das ist tatsächlich alles)

## Readiness Probe

Was ist das Standardverhalten (Es muss ein Prozess mit der id

Wann brauche ich die start

## Kubernetes - Wartung / Debugging

### Create Network-Connection to pod

#### Situation

```
Managed Cluster und ich kann nicht auf einzelne Nodes per ssh zugreifen  
Managed Cluster and i am not able to access nodes per ssh.
```

#### Behelf: Eigenen Pod starten mit busybox // Bring your own pod

```
## laengere Version / longer version  
kubectl run podtest --rm -ti --image busybox -- /bin/sh
```

```
## kuerzere Version  
kubectl run podtest --rm -ti --image busybox
```

#### Example test connection

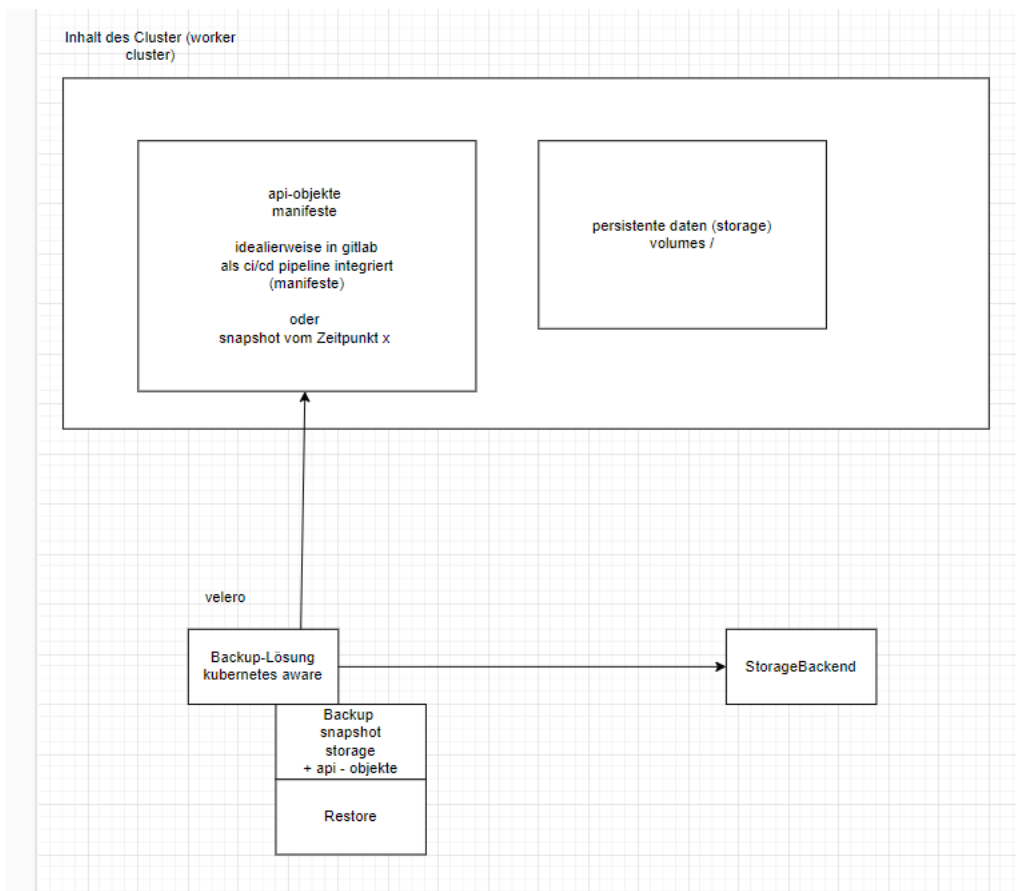
```
## wget (to copy)  
wget -O - http://10.244.0.99
```

```
## -O -> Output (grosses O (buchstabe))  
kubectl run podtest --rm -ti --image busybox -- /bin/sh  
/ # wget -O - http://10.244.0.99  
/ # exit
```

## Kubernetes Backup

### Backups mit Velero

#### Schaubild



#### Walkthrough in digitalocean

- <https://www.digitalocean.com/community/tutorials/how-to-back-up-and-restore-a-kubernetes-cluster-on-digitalocean-using-velero>

## Kubernetes Upgrade

### Upgrade von tanzu (Cluster API)

#### Step 1: Upgrade Tanzu Kubernetes Grid (Cluster Api)

- <https://docs.vmware.com/en/VMware-Tanzu-Kubernetes-Grid/1.6/vmware-tanzu-kubernetes-grid-16/GUID-upgrade-tkg-index.html>

#### Step 2: Upgrade Management Cluster

- <https://docs.vmware.com/en/VMware-Tanzu-Kubernetes-Grid/1.6/vmware-tanzu-kubernetes-grid-16/GUID-upgrade-tkg-management-cluster.html>

#### Step 3: Variante 1: Workload Cluster aktualisieren.

- <https://docs.vmware.com/en/VMware-Tanzu-Kubernetes-Grid/1.6/vmware-tanzu-kubernetes-grid-16/GUID-upgrade-tkg-workload-clusters.html>

#### Step 3: Variante 2: Neues Cluster hochziehen, ausrollen und altes abschalten

- <https://docs.vmware.com/en/VMware-Tanzu-Kubernetes-Grid/1.6/vmware-tanzu-kubernetes-grid-16/GUID-tanzu-k8s-clusters-index.html>

## Monitoring with Prometheus / Grafana

### Overview

#### What does it do ?

- It monitors your system by collecting data
- Data is pulled from your system by defined endpoints (http) from your cluster
- To provide data on your system, a lot of exporters are available, that
  - collect the data and provide it in Prometheus

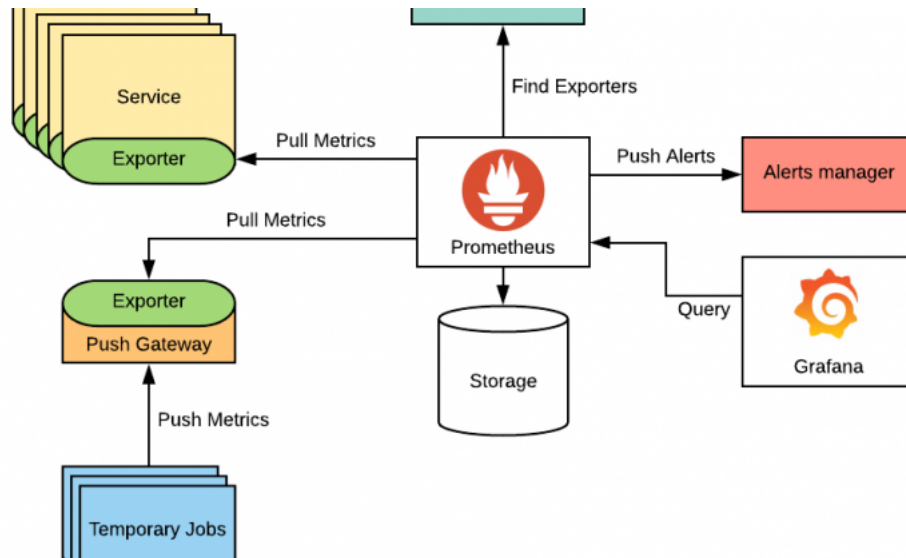
#### Technical

- Prometheus has a TDB (Time Series Database) and is good as storing time series with data
- Prometheus includes a local on-disk time series database, but also optionally integrates with remote storage systems.
- Prometheus's local time series database stores data in a custom, highly efficient format on local storage.
- Ref: <https://prometheus.io/docs/prometheus/latest/storage/>

## What are time series ?

- A time series is a sequence of data points that occur in successive order over some period of time.
- Beispiel:
  - Du willst die täglichen Schlusspreise für eine Aktie für ein Jahr dokumentieren
  - Damit willst Du weitere Analysen machen
  - Du würdest das Paar Datum/Preis dann in der Datumsreihenfolge sortieren und so ausgeben
  - Dies wäre eine "time series"

## Komponenten von Prometheus



Quelle: <https://www.devopsschool.com/>

## Prometheus Server

1. Retrieval (Sammeln)
  - Data Retrieval Worker
    - pull metrics data
2. Storage
  - Time Series Database (TDB)
    - stores metrics data
3. HTTP Server
  - Accepts PromQL - Queries (e.g. from Grafana)
    - accept queries

## Grafana ?

- Grafana wird meist verwendet um die grafische Auswertung zu machen.
- Mit Grafana kann ich einfach Dashboards verwenden
- Ich kann sehr leicht festlegen (Durch Data Sources), so meine Daten herkommen

## Setup prometheus/Grafana with helm

### Prerequisites

- Ubuntu 20.04 with running microk8s single cluster
- Works on any other cluster, but installing helm is different

### Prepare

```
## Be sure helm is installed on your client
## In our walkthrough, we will do it directly on 1 node,
## which is not recommended for Production
```

### Walkthrough

#### Step 1: install helm, if not there yet

```
snap install --classic helm
```

#### Step 2: Rollout prometheus/grafana stack in namespace prometheus



```
## add prometheus repo
helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
helm repo update

## install stack into new prometheus namespace
helm install -n prometheus --create-namespace prometheus prometheus-community/kube-prometheus-stack

## After installation look at the pods
## You should see 3 pods
kubectl --namespace prometheus get pods -l "release=prometheus"

## After a while it should be more pods
kubectl get all -n prometheus
```

### Step 3a Let's explain (der Prometheus - Server)

```
## 2 Stateful sets
kubectl get statefulsets -n prometheus
## output
## alertmanager-prometheus-kube-prometheus-alertmanager 1/1 5m14s
## prometheus-prometheus-kube-prometheus-prometheus. 1/1. 5m23s

## Moving part 1:
## prometheus-prometheus-kube-prometheus-prometheus
## That is the core prometheus server based on the main image

## Let's validate
## schauen wir mal in das File
kubectl get statefulset -n prometheus -o yaml > sts-prometheus-server.yaml

## Und vereinfacht (jetzt sehen wir direkt die beiden verwendeten images)
## 1) prometheus - server
## 2) der dazugehörige config-reloader als Side-Car
kubectl get sts -n prometheus prometheus-prometheus-kube-prometheus-prometheus -o
jsonpath='{.spec.template.spec.containers[*].image}'

## Aber wer managed den server -> managed-by -> kubernetes-operator
kubectl get sts -n prometheus prometheus-prometheus-kube-prometheus-prometheus -o jsonpath='{.spec.template.metadata.labels}' | jq
.

## Wir der sts von helm erstellt ?
## NEIN ;o)
## show us all the template that helm generate to apply them to kube-api-server
helm template prometheus prometheus-community/kube-prometheus-stack > all-prometheus.yaml
## NOPE -> none
cat all-prometheus.yaml | grep -i kind: | grep -i stateful

## secrets -> configuration von prometheus
## wenn ein eigenschaft Punkte hat, z.B. prometheus.yaml.gz
##
## {"prometheus.yaml.gz":"H4s
## dann muss man escapen, um darauf zuzugreifen -> aus . wird \.
kubectl get -n prometheus secrets prometheus-prometheus-kube-prometheus-prometheus -o jsonpath='{.data.prometheus\.yaml\.gz}' |
base64 -d | gzip -d -
```

### Step 3b: Prometheus Operator und Admission Controller -> Hook

```
## The Prometheus Operator for Kubernetes
## provides easy monitoring definitions
## for Kubernetes services and deployment and management of Prometheus instances.

## But how are they created
## After installation new resource-type are introduced
cat all-prometheus.yaml | grep ^kind: | grep -e 'Prometheus' -e 'ServiceM' | uniq
kind: Prometheus
kind: PrometheusRule
kind: ServiceMonitor
```

### Step 3c: How are the StatefulSets created

```
## New custom resource definitions are created
## The Prometheus custom resource definition (CRD) declaratively defines a desired Prometheus setup to run in a Kubernetes
cluster. It provides options to # configure replication, persistent storage, and Alertmanagers to which the deployed Prometheus
instances send alerts to.
```

```
## For each Prometheus resource, the Operator deploys a properly configured StatefulSet in the same namespace. The Prometheus Pods
are configured to mount # ca Secret called <prometheus-name> containing the configuration for Prometheus.
## https://github.com/prometheus-community/helm-charts/blob/main/charts/kube-prometheus-stack/crds/crd-prometheuses.yaml
```

### Step 3d: How are PrometheusRules created

```
## PrometheusRule are manipulated by the MutationHook when they enter the AdmissionController
## The AdmissionController is used after proper authentication in the kube-api-server
```

```
cat all-prometheus.yaml | grep 'Mutating' -B1 -A32
```

```
## Output
## Ref: https://kubernetes.io/docs/reference/access-authn-authz/extensible-admission-controllers/
apiVersion: admissionregistration.k8s.io/v1
kind: MutatingWebhookConfiguration
metadata:
  name: prometheus-kube-prometheus-admission
  labels:
    app: kube-prometheus-stack-admission
    app.kubernetes.io/managed-by: Helm
    app.kubernetes.io/instance: prometheus
    app.kubernetes.io/version: "35.4.2"
    app.kubernetes.io/part-of: kube-prometheus-stack
    chart: kube-prometheus-stack-35.4.2
    release: "prometheus"
    heritage: "Helm"
webhooks:
- name: prometheusrulemutate.monitoring.coreos.com
  failurePolicy: Ignore
  rules:
    - apiGroups:
        - monitoring.coreos.com
      apiVersions:
        - "*"
      resources:
        - prometheusrules
      operations:
        - CREATE
        - UPDATE
  clientConfig:
    service:
      namespace: prometheus
      name: prometheus-kube-prometheus-operator
      path: /admission-prometheusrules/mutate
    admissionReviewVersions: ["v1", "v1beta1"]
    sideEffects: None
```

### Step 4: Let's look into Deployments

```
kubectl -n prometheus get deploy
```

- What do they do

### Step 5: Let's look into DaemonSets

```
kubectl -n prometheus get ds
## node-exporter runs on every node
## connects to server, collects data and exports it
## so it is available for prometheus at the endpoint
```

### Helm -> prometheus stack -> What does it do

- Sets up Monitoring Stack
- Configuration for your K8s cluster
  - Worker Nodes monitored
  - K8s components (pods a.s.o) are monitored

### Where does configuration come from ?

```
## roundabout 31 configmaps
kubectl -n prometheus get configmaps

## also you have secrets (Grafana, Prometheus, Operator)
kubectl -n prometheus get secrets
```

### CRD's were created

```
## custom resource definitions
kubectl -n prometheus crd
## Sehr lang !
kubectl -n prometheus get crd/prometheuses.monitoring.coreos.com -o yaml
```

**Look into the pods to see the image used, how configuration is mounted**

```
kubectl -n prometheus get sts
kubectl -n prometheus describe sts/prometheus-prometheus-kube-prometheus-prometheus > prom.yml
kubectl -n prometheus describe sts/alertmanager-prometheus-kube-prometheus-alertmanager > alert.yml

kubectl -n prometheus get deploy
kubectl -n prometheus describe deploy/prometheus-kube-prometheus-operator > operator.yml

## ---> das SECRET erstellt der Kubernetes Operator für uns !
## First prom.yml
##. Mounts:
##   /etc/prometheus/config from config (rw)
##   -> What endpoints to scrape
## comes from:
kubectl get -n prometheus secrets prometheus-prometheus-kube-prometheus-prometheus -o jsonpath='{.data.prometheus\.yaml\.gz}' |
base64 -d | gunzip > config-prom.yml
## vi config-prom.yml
## Look into the scrape_configs
```

**Connect to grafana**

```
## wie ist der port 3000
kubectl logs prometheus-grafana-776fb976f7-w9nnp grafana
## hier nochmal port und auch, wie das secret heisst
kubectl describe pods prometheus-grafana-776fb976f7-w9nnp | less

## user / pass ?
kubectl get -n prometheus secrets prometheus-grafana -o jsonpath='{.data.admin-password}' | base64 -d
kubectl get -n prometheus secrets prometheus-grafana -o jsonpath='{.data.admin-user}' | base64 -d

## localhost:3000 erreichbarkeit starten -- im Vordergrund
kubectl port-forward deploy/prometheus-grafana 3000
## if on remote - system do a ssh-tunnel
## ssh -L 3000:127.0.0.1:3000 user@remote-ip

## letzte Schritt: browser aufrufen: http://localhost:3000
```

**Reference:**

- Techworld with Nana: <https://www.youtube.com/watch?v=QoDqxm7yblc>

**exporters mongodb**

**prometheus - export**

- <https://github.com/prometheus-community/helm-charts/tree/main/charts/prometheus-mongodb-exporter>

**Step 1: mongodb - deployment in mongodb namespace**

```
## vi mongo-db-deploy.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mongodb-deployment
  labels:
    app: mongodb
spec:
  replicas: 1
  selector:
    matchLabels:
      app: mongodb
  template:
    metadata:
      labels:
        app: mongodb
    spec:
      containers:
        - name: mongodb
          image: mongo
          ports:
```

```

- containerPort: 27017
---
apiVersion: v1
kind: Service
metadata:
  name: mongodb-service
spec:
  selector:
    app: mongodb
  ports:
    - protocol: TCP
      port: 27017
      targetPort: 27017

```

```
kubectl apply -f mongo-db-deploy.yml
```

## Step 2: Install prometheus - mongodb - export

```

helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
helm repo update
helm show values prometheus-community/prometheus-mongodb-exporter > values.yml

```

```

## adjust so it looks like so:
vi values.yml
## [mongodb[+srv]:/] [user:pass@]host1[:port1] [,host2[:port2],...][/database] [options]
## mongodb-service is the service name
mongodb:
  uri: "mongodb://mongodb-service:27017"

serviceMonitor:
  additionalLabels:
    release: prometheus

```

```
helm install mongodb-exporter prometheus-community/prometheus-mongodb-exporter -f values.yml
```

## Step 3: Helm -> template -> What does it do ?

```
helm template mongodb-exporter prometheus-community/prometheus-mongodb-exporter -f values.yml
```

## Good Kubernetes Board for Grafana

- <https://github.com/dotdc/grafana-dashboards-kubernetes>
- <https://medium.com/@dotdc/a-set-of-modern-grafana-dashboards-for-kubernetes-4b989c72a4b2>

## Kubernetes Tipps & Tricks

### kubectl kubeconfig mergen

#### So funktioniert es auch bereits:

```

## hier werden mehrere kubeconfigs durchsucht
export KUBECONFIG=~/.kube/config:/path/cluster1:/path/cluster2

```

#### Jetzt alles in eine Datei

```

cd ~/.kube
kubectl config view --flatten > all-in-one-kubeconfig.yaml
mv config config.old
mv all-in-one-kubeconfig.yaml config

```

#### Contexts jeweils anzeigen

```

kubectl config
kubectl config use-context mycontext

```

#### Create configmap from file

```

## Creates a users.yaml configmap-manifest
kubectl create cm users --from-file=users.json --dry-run=client -o yaml > users.yaml

```

## Kubernetes Certificates (Control Plane) / Security

### vmware - cluster api

- <https://docs.vmware.com/en/VMware-Tanzu-Kubernetes-Grid/1.6/vmware-tanzu-kubernetes-grid-16/GUID-cluster-lifecycle-secrets.html>

## Pod Security Admission (PSA)

### Seit: 1.2.22 Pod Security Admission

- 1.2.22 - Alpha - D.h. ist noch nicht aktiviert und muss als Feature Gate aktiviert (Kind)
- 1.2.23 - Beta -> d.h. evtl. aktiviert

### Vorgefertigte Regelwerke

- privileged - keinerlei Einschränkungen
- baseline - einige Einschränkungen
- restricted - sehr streng
- Reference: <https://kubernetes.io/docs/concepts/security/pod-security-standards/>

### Praktisches Beispiel für Version ab 1.2.23 - Problemstellung

```
mkdir -p manifests
cd manifests
mkdir psa
cd psa
nano 01-ns.yml
```

```
## Schritt 1: Namespace anlegen
## vi 01-ns.yml

apiVersion: v1
kind: Namespace
metadata:
  name: test-ns1
  labels:
    # soft version - running but showing complaints
    # pod-security.kubernetes.io/enforce: baseline
    pod-security.kubernetes.io/enforce: restricted
    pod-security.kubernetes.io/audit: restricted
    pod-security.kubernetes.io/warn: restricted
```

```
kubectl apply -f 01-ns.yml
```

```
## Schritt 2: Testen mit nginx - pod
## vi 02-nginx.yml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: test-ns1
spec:
  containers:
    - image: nginx
      name: nginx
      ports:
        - containerPort: 80
```

```
## a lot of warnings will come up
## because this image runs as root !! (by default)
kubectl apply -f 02-nginx.yml
```

```
## Schritt 3:
## Anpassen der Sicherheitseinstellung (Phase1) im Container
```

```
## vi 02-nginx.yml

apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: test-ns1
spec:
  containers:
    - image: nginx
      name: nginx
      ports:
        - containerPort: 80
      securityContext:
```

```
seccompProfile:
  type: RuntimeDefault
```

```
kubectl delete -f 02-nginx.yml
kubectl apply -f 02-nginx.yml
kubectl -n test-ns1 get pods
```

```
## Schritt 4:
## Weitere Anpassung runAsNotRoot
## vi 02-nginx.yml
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: test-ns<tl>
spec:
  containers:
    - image: nginx
      name: nginx
      ports:
        - containerPort: 80
      securityContext:
        seccompProfile:
          type: RuntimeDefault
        runAsNonRoot: true
```

```
## pod kann erstellt werden, wird aber nicht gestartet
kubectl delete -f 02-nginx.yml
kubectl apply -f 02-nginx.yml
kubectl -n test-ns1 get pods
kubectl -n test-ns1 describe pods nginx
```

```
## Schritt 4:
## Anpassen der Sicherheitseinstellung (Phase1) im Container
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: test-ns1
spec:
  containers:
    - image: nginx
      name: nginx
      ports:
        - containerPort: 80
      securityContext:
        seccompProfile:
          type: RuntimeDefault
        runAsNonRoot: true
        allowPrivilegeEscalation: false
        capabilities:
          drop: ["ALL"]
```

```
kubectl delete -f 02-nginx.yml
kubectl apply -f 02-nginx.yml
kubectl -n test-ns1 get pods
```

### Praktisches Beispiel für Version ab 1.2.23 -Lösung - Container als NICHT-Root laufen lassen

- Wir müssen ein image, dass auch als NICHT-Root laufen kann
- .. oder selbst eines bauen (:o)) o bei nginx ist das bitnami/nginx

```
## vi 03-nginx-bitnami.yml
apiVersion: v1
kind: Pod
metadata:
  name: bitnami-nginx
  namespace: test-ns1
spec:
  containers:
    - image: bitnami/nginx
      name: bitnami-nginx
      ports:
        - containerPort: 80
      securityContext:
```

```
seccompProfile:
  type: RuntimeDefault
runAsNonRoot: true
```

```
## und er läuft als nicht root
kubectl apply -f 03_pod-bitnami.yml
kubectl -n test-ns1 get pods
```

## Pod Security Policy (PSP)

### General

- PodSecurity is an eine Rolle gebunden (clusterrole)
- Deprecated in 1.21 removed in 1.25
- From 1.25 on please use PSA (Pod Security Admission) instead

### Prerequisites

- We should have a running Cluster of 1.22/1.23

### Walkthrough

#### Step 1: Create Digitalocean microk8s 1-node - cluster, with this cloud-init-script

- cloud-init (ubuntu 20.04 LTS, 8 GB Ram)

```
#!/bin/bash

groupadd sshadmin
USERS="11trainingdo"
echo $USERS
for USER in $USERS
do
    echo "Adding user $USER"
    useradd -s /bin/bash --create-home $USER
    usermod -aG sshadmin $USER
    echo "$USER:deinsehrgeheimesspasswort" | chpasswd
done

## We can sudo with 11trainingdo
usermod -aG sudo 11trainingdo

## 20.04 and 22.04 this will be in the subfolder
if [ -f /etc/ssh/sshd_config.d/50-cloud-init.conf ]
then
    sed -i "s/PasswordAuthentication no/PasswordAuthentication yes/g" /etc/ssh/sshd_config.d/50-cloud-init.conf
fi

### both is needed
sed -i "s/PasswordAuthentication no/PasswordAuthentication yes/g" /etc/ssh/sshd_config

usermod -aG sshadmin root

## TBD - Delete AllowUsers Entries with sed
## otherwise we cannot login by group

echo "AllowGroups sshadmin" >> /etc/ssh/sshd_config
systemctl reload sshd

echo "Installing microk8s"
snap install --classic --channel=1.23/stable microk8s
microk8s enable dns rbac
echo "alias kubectl='microk8s kubectl'" >> /root/.bashrc
source ~/.bashrc
alias kubectl='microk8s kubectl'

## now we need to modify the setting of kube-api-server
## currently in 1.23 no other admission-plugins are activated
echo "--enable-admission-plugins=PodSecurityPolicy" >> /var/snap/microk8s/current/args/kube-apiserver
microk8s stop
microk8s start
```

#### Step 2:

```
## Setup .kube/config from content
microk8s config
```

#### Step 3

```

## rbac.yaml
## vi service-account.yml
apiVersion: v1
kind: ServiceAccount
metadata:
  name: training
  namespace: default
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: pods-clusterrole
rules:
- apiGroups: [""] # "" indicates the core API group
  resources: ["pods"]
  verbs: ["get", "watch", "list", "create"]
- apiGroups: [""] # "" indicates the core API group
  resources: ["events"]
  verbs: ["get", "list"]
---
## vi rb-training-ns-default-pods.yml
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: rolebinding-ns-default-pods
  namespace: default
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: pods-clusterrole
subjects:
- kind: ServiceAccount
  name: training
  namespace: default

```

#### Step 4: Secret aus secrets rauskopiert

```

kubectl get secrets | grep training-token
TOKEN=$(kubectl get secrets training-token-xyz -o jsonpath='{.data.token}' | base64 -d)
## z.B. TOKEN=$(kubectl get secrets training-token-kjl5m -o jsonpath='{.data.token}' | base64 -d)

echo $TOKEN
kubectl config set-context training-ctx --cluster microk8s-cluster --user training
kubectl config set-credentials training --token=$TOKEN

```

#### Step 5: Apply yaml-manifests for psp - stuff (as admin)

```

## vi setup.yaml
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  name: norootcontainers
spec:
  allowPrivilegeEscalation: false
  allowedHostPaths:
  - pathPrefix: /dev/null
    readOnly: true
  fsGroup:
    rule: RunAsAny
  hostPorts:
  - max: 65535
    min: 0
  runAsUser:
    rule: MustRunAsNonRoot
  seLinux:
    rule: RunAsAny
  supplementalGroups:
    rule: RunAsAny
  volumes:
  - '*'
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: norootcontainers-ppsp-role
rules:

```



```

- apiGroups:
  - policy
  resourceNames:
  - norootcontainers
  resources:
  - podsecuritypolicies
  verbs:
  - use
---
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: norootcontainers-psp-role:training
  namespace: default
roleRef:
  kind: ClusterRole
  name: norootcontainers-psp-role
  apiGroup: rbac.authorization.k8s.io
subjects:
- kind: ServiceAccount
  name: training
  namespace: default

```

### Step 5: Change to training-ctx and apply

```
kubectl config use-context training-ctx
```

```

## vi demopod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: demopod
spec:
  containers:
  - name: demopod
    image: nginx

```

```

kubectl apply -f demopod.yaml
kubectl get pods ## expecting
kubectl describe pods demopod

```

### Reference

- <https://docs.mirantis.com/mke/3.4/ops/deploy-apps-k8s/pod-security-policies/psp-examples.html>

## Kubernetes Network / Firewall

### Calico/Cilium - nginx example NetworkPolicy

```

## Schritt 1:
kubectl create ns policy-demo
kubectl create deployment --namespace=policy-demo nginx --image=nginx:1.21
kubectl expose --namespace=policy-demo deployment nginx --port=80
## lassen einen 2. pod laufen mit dem auf den nginx zugreifen
kubectl run --namespace=policy-demo access --rm -it --image busybox

```

```

## innerhalb der shell
## Verbindung möglich
wget -q nginx -O -

```

```

## Schritt 2: Policy festlegen, dass kein Ingress-Traffic erlaubt
## in diesem namespace: policy-demo
kubectl create -f - <<EOF
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: default-deny
  namespace: policy-demo
spec:
  podSelector:
    matchLabels: {}
EOF
## lassen einen 2. pod laufen mit dem auf den nginx zugreifen
kubectl run --namespace=policy-demo access --rm -ti --image busybox

```

```
## innerhalb der shell
## keine Verbindung mehr möglich, weil policy greift
wget -q nginx -O -
```

```
## Schritt 3: Zugriff erlauben von pods mit dem Label run=access
kubectl create -f - <<EOF
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: access-nginx
  namespace: policy-demo
spec:
  podSelector:
    matchLabels:
      app: nginx
  ingress:
    - from:
      - podSelector:
          matchLabels:
            run: access
EOF

## lassen einen 2. pod laufen mit dem auf den nginx zugreifen
## pod hat durch run -> access automatisch das label run:access zugewiesen
kubectl run --namespace=policy-demo access --rm -ti --image busybox
```

```
## innerhalb der shell
wget -q nginx -O -
```

```
kubectl run --namespace=policy-demo no-access --rm -ti --image busybox
```

```
## in der shell
wget -q nginx -O -
```

```
kubectl delete ns policy-demo
```

#### Ref:

- <https://projectcalico.docs.tigera.io/security/tutorials/kubernetes-policy-basic>

#### Egress / Ingress Examples with Exercise

##### Links

- <https://github.com/ahmetb/kubernetes-network-policy-recipes>
- <https://k8s-examples.container-solutions.com/examples/NetworkPolicy/NetworkPolicy.html>

#### Example with http (Cilium !!)

```
apiVersion: "cilium.io/v2"
kind: CiliumNetworkPolicy
description: "L7 policy to restrict access to specific HTTP call"
metadata:
  name: "rule1"
spec:
  endpointSelector:
    matchLabels:
      type: l7-test
  ingress:
    - fromEndpoints:
      - matchLabels:
          org: client-pod
      toPorts:
        - ports:
            - port: "8080"
              protocol: TCP
        rules:
          http:
            - method: "GET"
              path: "/discount"
```

#### Downside egress (NetworkPolicy - not ciliumnetworkpolicy)

- No valid api for anything other than IP's and/or Ports
- If you want more, you have to use CNI-Plugin specific, e.g.

#### Example egress with ip's

```
## Allow traffic of all pods having the label role:app
## egress only to a specific ip and port
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: test-network-policy
  namespace: default
spec:
  podSelector:
    matchLabels:
      role: app
  policyTypes:
  - Egress
  egress:
  - to:
    - ipBlock:
        cidr: 10.10.0.0/16
      ports:
      - protocol: TCP
        port: 5432
```

### Example Advanced Egress (cni-plugin specific)

#### Cilium (Exercise)

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-static-web
  labels:
    webserver: nginx
spec:
  containers:
  - name: web
    image: nginx
```

```
apiVersion: cilium.io/v2
kind: CiliumNetworkPolicy
metadata:
  name: "fqdn-pprof"
  # namespace: msp
spec:
  endpointSelector:
    matchLabels:
      webserver: nginx
  egress:
  - toFQDNs:
    - matchPattern: '*.google.com'
  - toPorts:
    - ports:
      - port: "53"
        protocol: ANY
      rules:
        dns:
        - matchPattern: '*'
```

```
kubectl apply -f .
kubectl exec -it nginx-static-web -- bash
```

```
## im pod
## does work
curl -I https://www.google.com
## does not work
curl -I https://www.google.de
## does not work
curl -I https://www.heise.de
```

#### Calico

- Only Calico enterprise
  - Calico Enterprise extends Calico's policy model so that domain names (FQDN / DNS) can be used to allow access from a pod or set of pods (via label selector) to external resources outside of your cluster.
  - <https://projectcalico.docs.tigera.io/security/calico-enterprise/egress-access-controls>

#### Using istio as mesh (e.g. with cilium/calico )

##### Installation of sidecar in calico

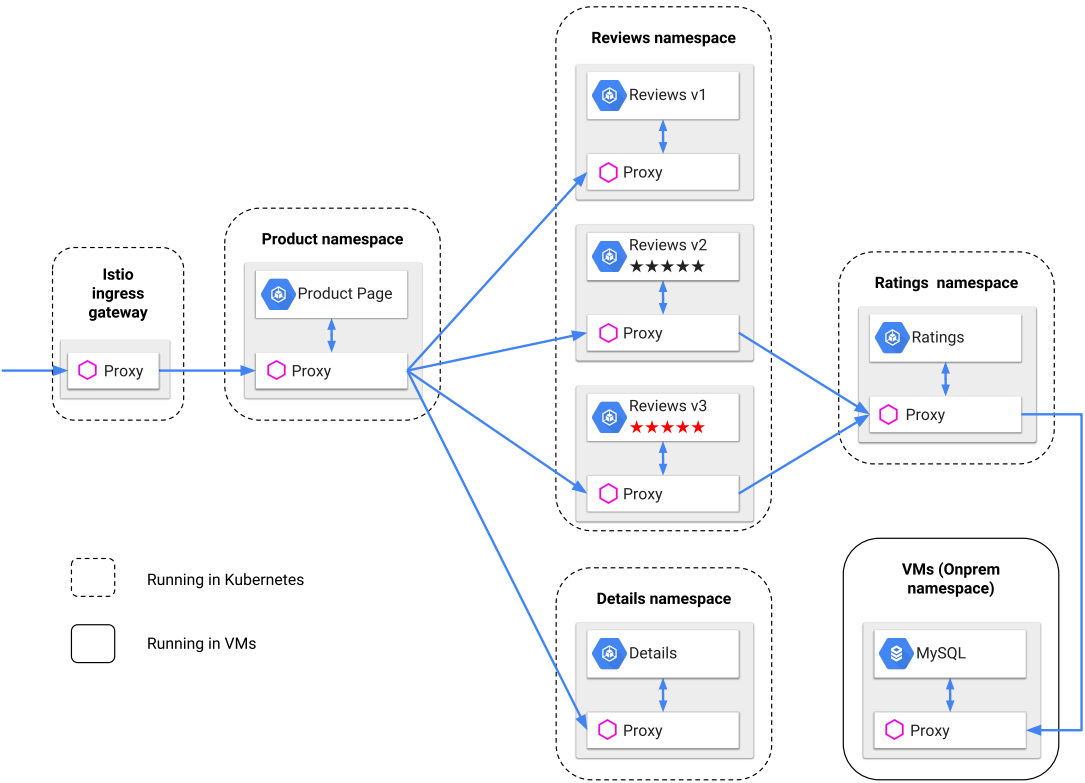
- <https://projectcalico.docs.tigera.io/getting-started/kubernetes/hardway/istio-integration>

Example

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: test-network-policy
  namespace: default
spec:
  podSelector:
    matchLabels:
      role: app
  policyTypes:
    - Egress
  egress:
    - to:
      - ipBlock:
          cidr: 10.10.0.0/16
    ports:
      - protocol: TCP
        port: 5432
```

Mesh / istio

Schaubild



Istio

```
## Visualization
## with kiali (included in istio)
https://istio.io/latest/docs/tasks/observability/kiali/kiali-graph.png

## Example
## https://istio.io/latest/docs/examples/bookinfo/
The sidecars are injected in all pods within the namespace by labeling the namespace like so:
kubectl label namespace default istio-injection=enabled

## Gateway (like Ingress in vanilla Kubernetes)
kubectl label namespace default istio-injection=enabled
```

istio tls

- <https://istio.io/latest/docs/ops/configuration/traffic-management/tls-configuration/>

## istio - the next generation without sidecar

- <https://istio.io/latest/blog/2022/introducing-ambient-mesh/>

## Kubernetes Probes (Liveness and Readiness)

### Übung Liveness-Probe

#### Übung 1: Liveness (command)

What does it do ?

```
* At the beginning pod is ready (first 30 seconds)
* Check will be done after 5 seconds of pod being startet
* Check will be done periodically every 5 minutes and will check
  * for /tmp/healthy
  * if file is there will return: 0
  * if file is not there will return: 1
* After 30 seconds container will be killed
* After 35 seconds container will be restarted
```

```
cd
mkdir -p manifests/probes
cd manifests/probes
## vi 01-pod-liveness-command.yml
```

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    test: liveness
  name: liveness-exec
spec:
  containers:
    - name: liveness
      image: busybox
      args:
        - /bin/sh
        - -c
        - touch /tmp/healthy; sleep 30; rm -f /tmp/healthy; sleep 600
      livenessProbe:
        exec:
          command:
            - cat
            - /tmp/healthy
        initialDelaySeconds: 5
        periodSeconds: 5
```

```
## apply and test
kubectl apply -f 01-pod-liveness-command.yml
kubectl describe -l test=liveness pods
sleep 30
kubectl describe -l test=liveness pods
sleep 5
kubectl describe -l test=liveness pods
```

```
## cleanup
kubectl delete -f 01-pod-liveness-command.yml
```

#### Übung 2: Liveness Probe (HTTP)

```
## Step 0: Understanding Prerequisite:
This is how this image works:
## after 10 seconds it returns code 500
http.HandleFunc("/healthz", func(w http.ResponseWriter, r *http.Request) {
    duration := time.Now().Sub(started)
    if duration.Seconds() > 10 {
        w.WriteHeader(500)
        w.Write([]byte(fmt.Sprintf("error: %v", duration.Seconds())))
    } else {
        w.WriteHeader(200)
        w.Write([]byte("ok"))
    }
})
```

```
## Step 1: Pod - manifest
## vi 02-pod-liveness-http.yml
## status-code >=200 and < 400 o.k.
## else failure
apiVersion: v1
kind: Pod
metadata:
  labels:
    test: liveness
    name: liveness-http
spec:
  containers:
  - name: liveness
    image: k8s.gcr.io/liveness
    args:
    - /server
    livenessProbe:
      httpGet:
        path: /healthz
        port: 8080
        httpHeaders:
        - name: Custom-Header
          value: Awesome
      initialDelaySeconds: 3
      periodSeconds: 3
```

```
## Step 2: apply and test
kubectl apply -f 02-pod-liveness-http.yml
## after 10 seconds port should have been started
sleep 10
kubectl describe pod liveness-http
```

#### Reference:

- <https://kubernetes.io/docs/tasks/configure-pod-container/configure-liveness-readiness-startup-probes/>

### Übung Liveness http aus nginx

#### Funktionsweise Readiness-Probe vs. Liveness-Probe

##### Why / Howto /

- Readiness checks, if container is ready and if it's not READY
  - SENDS NO TRAFFIC to the container

#### Difference to LiveNess

- They are configured exactly the same, but use another keyword
  - readinessProbe instead of livenessProbe

#### Example

```
readinessProbe:
  exec:
    command:
    - cat
    - /tmp/healthy
  initialDelaySeconds: 5
  periodSeconds: 5
```

#### Reference

- <https://kubernetes.io/docs/tasks/configure-pod-container/configure-liveness-readiness-startup-probes/#define-readiness-probes>

### Manueller Check readyz endpoint kubernetes api server aus pod

#### Walkthrough

```
kubectl run -it --rm podtester --image=busybox
## im pod
## um zu sehen mit welchem Port wir uns verbinden können
env | grep -i kubernetes
## kubernetes liegt als service vor
wget -O - https://kubernetes:443/readyz?verbose
```

#### Reference:

- <https://kubernetes.io/docs/reference/using-api/health-checks/>

### Kubernetes QoS / Limits / Requests

## Quality of Service - evict pods

### Die Class wird auf Basis der Limits und Requests der Container vergeben

#### Request

Request: Definiert wieviel ein Container mindestens braucht (CPU,memory)

#### Limit

Limit: Definiert, was ein Container maximal braucht.

#### Wo ?

```
in spec.containers.resources
kubectl explain pod.spec.containers.resources
```

#### Art der Typen:

- Guaranteed
- Burstable
- BestEffort

#### Guaranteed

```
Type: Guaranteed:
https://kubernetes.io/docs/tasks/configure-pod-container/quality-service-pod/#create-a-pod-that-gets-assigned-a-qos-class-of-guaranteed

set when limit equals request
(request: das braucht er,
limit: das braucht er maximal)

Garantied ist die höchste Stufe und diese werden bei fehlenden Ressourcen
als letztes "evicted"

apiVersion: v1

kind: Pod
metadata:
  name: qos-demo
  namespace: qos-example
spec:
  containers:
  - name: qos-demo-ctr
    image: nginx
    resources:
      limits:
        memory: "200Mi"
        cpu: "700m"

      requests:
        memory: "200Mi"
        cpu: "700m"
```

#### Referenz

- <https://home.robusta.dev/blog/kubernetes-memory-limit>

#### Tools to Identify LimitRange and Requests

#### VPA (Vertical Pod Autoscaler) / goldilocks

```
## Please only repo updateMode: "off" will do this
## Do not use automatic adjustment
Example VPA configuration
apiVersion: autoscaling.k8s.io/v1
kind: VerticalPodAutoscaler
metadata:
  name: my-app-vpa
spec:
  targetRef:
    apiVersion: "apps/v1"
    kind: Deployment
    name: my-app
  updatePolicy:
    updateMode: "off"
```

- goldilocks will now make visible instead of kubectl describe vpa
- <https://github.com/FairwindsOps/goldilocks>
- als Basis: <https://github.com/kubernetes/autoscaler/>
- <https://www.fairwinds.com/goldilocks>

## Kubernetes Autoscaling

### Autoscaling Pods/Deployments

**Example: newest version with autoscaling/v2 used to be hpa/v1**

#### Prerequisites

- Metrics-Server needs to be running

```
## Test with
kubectl top pods
```

```
## Install
kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
## after that it will be available in kube-system namespace as pod
kubectl -n kube-system get pods | grep -i metrics
```

#### Step 1: deploy app

```
cd
mkdir -p manifests
cd manifests
mkdir hpa
cd hpa
vi 01-deploy.yaml
```

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello
spec:
  replicas: 3
  selector:
    matchLabels:
      app: hello
  template:
    metadata:
      labels:
        app: hello
    spec:
      containers:
        - name: hello
          image: k8s.gcr.io/hpa-example
          resources:
            requests:
              cpu: 100m
---
kind: Service
apiVersion: v1
metadata:
  name: hello
spec:
  selector:
    app: hello
  ports:
    - port: 80
      targetPort: 80
---
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: hello
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: hello
  minReplicas: 2
  maxReplicas: 20
  metrics:
    - type: Resource
```



```

resource:
  name: cpu
  target:
    type: Utilization
    averageUtilization: 80

```

## Step 2: Load Generator

```
vi 02-loadgenerator.yml
```

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: load-generator
  labels:
    app: load-generator
spec:
  replicas: 100
  selector:
    matchLabels:
      app: load-generator
  template:
    metadata:
      name: load-generator
      labels:
        app: load-generator
    spec:
      containers:
        - name: load-generator
          image: busybox
          command:
            - /bin/sh
            - -c
            - "while true; do wget -q -O- http://hello.default.svc.cluster.local; done"

```

## Downscaling

- Downscaling will happen after 5 minutes o

```

## Adjust down to 1 minute
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: hello
spec:
  # change to 60 secs here
  behavior:
    scaleDown:
      stabilizationWindowSeconds: 60
  # end of behaviour change
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: hello
  minReplicas: 2
  maxReplicas: 20
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 80

```

For scaling down the stabilization window is 300 seconds (or the value of the --horizontal-pod-autoscaler-downscale-stabilization flag if provided)

## Prevent Downscaling

```

## Adjust down to 1 minute
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: hello
spec:

```

```
# change to 60 secs here
behavior:
  scaleDown:
    selectPolicy: Disabled
# end of behaviour change
scaleTargetRef:
  apiVersion: apps/v1
  kind: Deployment
  name: hello
minReplicas: 2
maxReplicas: 20
metrics:
- type: Resource
  resource:
    name: cpu
    target:
      type: Utilization
      averageUtilization: 80
```

## Reference

- <https://docs.digitalocean.com/tutorials/cluster-autoscaling-ca-hpa/>
- <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale-walkthrough/#autoscaling-on-more-specific-metrics>
- <https://medium.com/expedia-group-tech/autoscaling-in-kubernetes-why-doesnt-the-horizontal-pod-autoscaler-work-for-me-5f0094694054>

## Kubernetes Deployment Scenarios

### Deployment green/blue,canary,rolling update

#### Canary Deployment

A small group of the user base will see the new application  
(e.g. 1000 out of 100.000), all the others will still see the old version

From: a canary was used to test if the air was good in the mine  
(like a test balloon)

#### Blue / Green Deployment

The current version is the Blue one  
The new version is the Green one

New Version (GREEN) will be tested and if it works  
the traffic will be switch completey to the new version (GREEN)

Old version can either be deleted or will function as fallback

#### A/B Deployment/Testing

2 Different versions are online, e.g. to test a new design / new feature  
You can configure the weight (how much traffic to one or the other)  
by the number of pods

#### Example Calculation

e.g. Deployment1: 10 pods  
Deployment2: 5 pods

Both have a common label,  
The service will access them through this label

#### Service Blue/Green

##### Step 1: Deployment + Service

```
## vi blue.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-version-blue
spec:
  selector:
    matchLabels:
      version: blue
  replicas: 10 # tells deployment to run 2 pods matching the template
```

```

template:
  metadata:
    labels:
      app: nginx
      version: blue
  spec:
    containers:
      - name: nginx
        image: nginx:1.21
        ports:
          - containerPort: 80

```

```

## vi green.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-version-green
spec:
  selector:
    matchLabels:
      version: green
  replicas: 1 # tells deployment to run 2 pods matching the template
  template:
    metadata:
      labels:
        app: nginx
        version: green
    spec:
      containers:
        - name: nginx
          image: nginx:1.22
          ports:
            - containerPort: 80

```

```

## svc.yml
apiVersion: v1
kind: Service
metadata:
  name: svc-nginx
spec:
  ports:
    - port: 80
      protocol: TCP
  selector:
    app: nginx

```

## Step 2: Ingress

```

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: ingress-config
  annotations:
    ingress.kubernetes.io/rewrite-target: /
    # with the ingress controller from helm, you need to set an annotation
    # old version useClassName instead
    # otherwise it does not know, which controller to use
    # kubernetes.io/ingress.class: nginx
spec:
  ingressClassName: nginx
  rules:
    - host: "app.labl.t3isp.de"
      http:
        paths:
          - path: /
            pathType: Prefix
            backend:
              service:
                name: svc-nginx
                port:
                  number: 80

```

```
kubectl apply -f .
```

## Praxis-Übung A/B Deployment

## Walkthrough

```
cd
cd manifests
mkdir ab
cd ab
```

```
## vi 01-cm-version1.yml
apiVersion: v1
kind: ConfigMap
metadata:
  name: nginx-version-1
data:
  index.html: |
    <html>
    <h1>Welcome to Version 1</h1>
    </br>
    <h1>Hi! This is a configmap Index file Version 1 </h1>
    </html>
```

```
## vi 02-deployment-v1.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deploy-v1
spec:
  selector:
    matchLabels:
      version: v1
  replicas: 2
  template:
    metadata:
      labels:
        app: nginx
        version: v1
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
          volumeMounts:
            - name: nginx-index-file
              mountPath: /usr/share/nginx/html/
      volumes:
        - name: nginx-index-file
          configMap:
            name: nginx-version-1
```

```
## vi 03-cm-version2.yml
apiVersion: v1
kind: ConfigMap
metadata:
  name: nginx-version-2
data:
  index.html: |
    <html>
    <h1>Welcome to Version 2</h1>
    </br>
    <h1>Hi! This is a configmap Index file Version 2 </h1>
    </html>
```

```
## vi 04-deployment-v2.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deploy-v2
spec:
  selector:
    matchLabels:
      version: v2
  replicas: 2
  template:
    metadata:
      labels:
        app: nginx
```

```
    version: v2
spec:
  containers:
  - name: nginx
    image: nginx:latest
    ports:
    - containerPort: 80
    volumeMounts:
    - name: nginx-index-file
      mountPath: /usr/share/nginx/html/
  volumes:
  - name: nginx-index-file
    configMap:
      name: nginx-version-2
```


```
## vi 05-svc.yml
apiVersion: v1
kind: Service
metadata:
  name: my-nginx
  labels:
    svc: nginx
spec:
  type: NodePort
  ports:
  - port: 80
    protocol: TCP
  selector:
    app: nginx
```

```
kubectl apply -f .
## get external ip
kubectl get nodes -o wide
## get port
kubectl get svc my-nginx -o wide
## test it with curl apply it multiple time (at least ten times)
curl <external-ip>:<node-port>
```

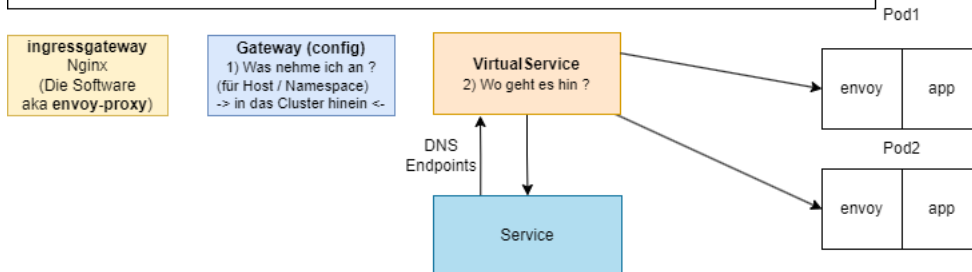
## Kubernetes Istio

### Istio vs. Ingress Überblick

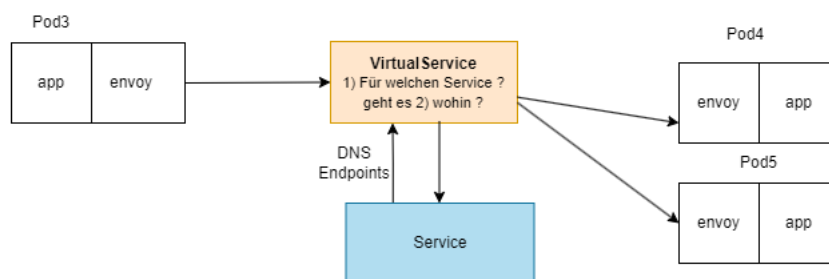
## Klassisch Ingress (Kubernetes)



## Istio (Traffic in das Cluster hinein)



## Istio (Innerhalb des Clusters)



## Istio installieren und Addons bereitstellen

### On the client (where you also use kubectl)

#### Steps 1: Download install and run

```
## as tlnx - user
## find a decent where to run the installation
## not perfect, but better than to put it in home-folder
cd
mkdir -p manifests/istio
cd manifests/istio
```

```
## now download the install an run the shell
curl -L https://istio.io/downloadIstio | sh -
```

#### Step 2: Run istioctl - commands (version-check, precheck and install)

```
## This istioctl will be under istio-1.20.2/bin
## but TRAINER has already installed it under /usr/bin/istioctl
## So we can use that one !!
```

```
## cd istio-1.20.2/bin
istioctl version
istioctl x precheck
istioctl install --set profile=demo -y
```

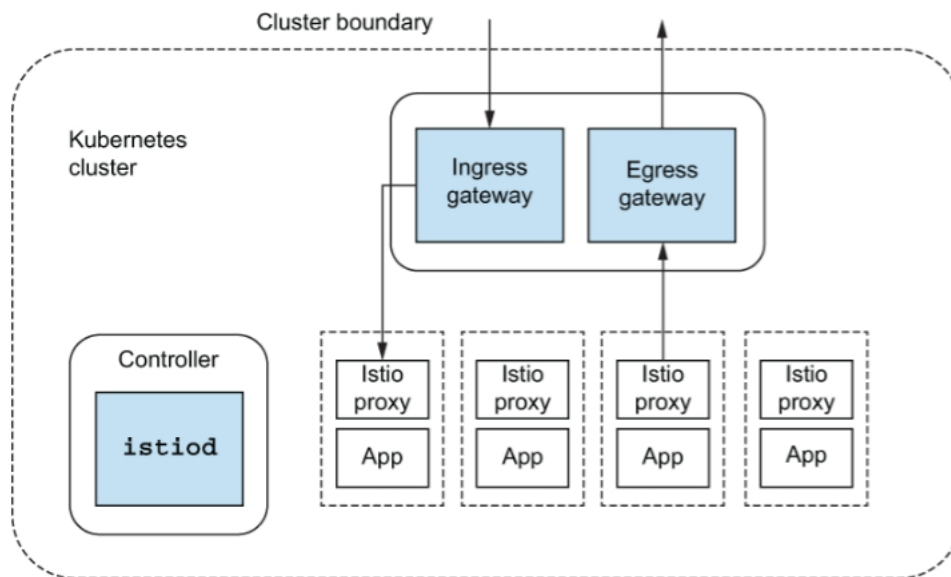
#### Step 3: Install the addons

```
## Install Add-Ons
kubectl apply -f istio-1.20.2/samples/addons/
```

#### Step 4: Check if all the corresponding container (from istio and addons) are running


```
kubectl -n istio-system get pods
```

## Istion Überblick - egress und ingress - gateway



## Istio - Deployment of simple application

### Overview (what we want to do)



- Catalog Service is reachable through api

### Step 1: Vorbereitung - repo mit beispilen klonen

```
cd
git clone https://github.com/jmetzger/istio-exercises/
cd istio-exercises
```

### Step 2: Eigenen Namespace erstellen

```
## Jeder Teilnehmer erstellt seinen eigenen Namespace
## z.B. istioapp-tlnx
## d.h. für Teilnehmer 5 (tln5) -> istioapp-tln5
kubectl create ns istioapp-tln5
## Context so einstellen, dass dieser namespace verwendet
kubectl config set-context --current --namespace istioapp-tln5
```

### Step 3: Anwendung untersuchen / istioctl kube-inject

- Ihr könnt unten direkt den Pfad nehmen, das ist einfacher ;o)

```
apiVersion: v1
kind: ServiceAccount
metadata:
```

```

  name: catalog
---
apiVersion: v1
kind: Service
metadata:
  labels:
    app: catalog
  name: catalog
spec:
  ports:
    - name: http
      port: 80
      protocol: TCP
      targetPort: 3000
  selector:
    app: catalog
---
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: catalog
    version: v1
  name: catalog
spec:
  replicas: 1
  selector:
    matchLabels:
      app: catalog
      version: v1
  template:
    metadata:
      labels:
        app: catalog
        version: v1
    spec:
      serviceAccountName: catalog
      containers:
        - env:
            - name: KUBERNETES_NAMESPACE
              valueFrom:
                fieldRef:
                  fieldPath: metadata.namespace
            image: istioinaction/catalog:latest
            imagePullPolicy: IfNotPresent
            name: catalog
            ports:
              - containerPort: 3000
                name: http
                protocol: TCP
            securityContext:
              privileged: false

```

```

## schauen wir uns das mal mit injection an
istioctl kube-inject -f services/catalog/kubernetes/catalog.yaml | less

```

#### Step 4: Automatische Injection einrichten.

```

## kubectl label namespace istioapp-tlnx istio-injection=enabled
## z.B
kubectl label namespace istioapp-tln1 istio-injection=enabled

```

#### Step 5: catalog ausrollen

```

kubectl apply -f services/catalog/kubernetes/catalog.yaml

```

```

## Prüfen, ob wirklich 2 container in einem pod laufen,
## dann funktioniert die Injection
## WORKS, Yeah !
kubectl get pods

```

#### Step 6: Wir wollen den Catalog jetzt erreichen

```

## do it from your namespace, e.g. tlnx
## z.B.
kubectl -n tln1 run -it --rm curly --image=curlimages/curl -- sh

```



```
## within shell of that pod
## catalog.yourappnamespace/items/1
curl http://catalog.istioapp-tln1/items/1
exit
```

### Step 7: Jetzt deployen wir die webapp

```
## Wir schauen uns das manifest für die webapp an
## und ändern die env-variablen CATALOG_SERVICE_HOST
## tlnx durch Eure Teilnehmernummer ersetzen
catalog.istioapp-tlnx
```

```
kubectl apply -f services/webapp/kubernetes/webapp.yaml
kubectl get pod
```

### Step 8: Verbindung zu webapp testen

```
## tlnx
## kubectl -n tlnx run -it --rm curly --image=curlimages/curl -- sh
## z.B.
kubectl -n tln5 run -it --rm curly --image=curlimages/curl -- sh
```

```
## Within shell connect to webapp
curl -s http://webapp.istioapp-tln1/api/catalog/items/1
exit
```

```
## Wir können es aber auch visualisieren
kubectl port-forward deploy/webapp 8001:8080
## z.B. Teilnehmer tln1 -> 8001:8080

## WICHTIG Jeder Teilnehmer sollte hier einen abweichenden Port nehmen
## Jetzt lokal noch einen Tunnel aufbauen
## s. Anleitung Putty
## Source Port: 8080 # das ist der auf dem Rechner
## Destination: localhost:8001
## Add
## Achtung -> danach noch Session speichern
```

```
## Jetzt im Browser http://localhost:8080
## aufrufen
```

### Step 9: Ingress - Gateway konfigurieren (ähnlich wie Ingress-Objekt)

```
## wir schauen uns das vorher mal an
```

```
## namespace - fähig, d.h. ein Gateway mit gleichem Namen pro Namespace möglich
cat ingress-virtualservice/ingress-gateway.yaml
## hier bitte bei Hosts hostname eintragen, der für t3isp.de verwendet, und zwar
## jeder Teilnehmer eine eigene Subdomain: z.B. jochen.istio.t3isp.de
kubectl apply -f ingress-virtualservice/ingress-gateway.yaml
```

### Step 10: Reach it from outside

```
## We need to find the loadbalancer IP
kubectl -n istio-system get svc
## in unserem Fall
146.190.177.12
## Das trägt Jochen dns t3isp.de ein.
```

```
## Wir können jetzt also das System von extern erreichen
## vomn client aus, oder direkt über den Browser
##curl -i 146.190.177.12/api/catalog/items/1
## Hier hostname statt ip eintragen
curl -i http://tlnx.istio.t3isp.de/api/catalog/items/1
```

```
## Wir können auch über istioctl direkt überprüfen, ob es einen Routen-Config gibt
istioctl proxy-config routes deploy/istio-ingressgateway.istio-system
```

```
## Falls das nicht funktioniert, können wir auch überprüfen ob ein gateway und ein virtualservice installiert wurde
kubectl get gateway
kubectl get virtualservice
## Kurzform des Services reicht, weil im gleichen namespace
## Wo soll es hingehen -> == -> Upstream
```

```
## route -> destination -> host -> webapp
kubectl get virtualservice -o yaml
```

```
### Wichtiger Hinweis, auf beiden Seiten ingressgateway und vor dem Pod des Dienstes Webapp
### Sitzt ein envoy-proxy und kann Telemetrie-Daten und Insight sammeln was zwischen den
### applicationen passiert -> das passiert über ein sidecar in jeder Applikation
```

```
### Wichtig: Das passiert alles ausserhalb der Applikation
### Nicht wie früher z.B. bei Netflix innerhalb z.B. für die Sprache Java
```

## Istio - Grafana Dashboard

### Status

- Wir haben bereits mit den Addons Grafana ausgerollt,
- Dieses wollen wir jetzt aktivieren

### Schritt 1: Dashboard aktivieren -> achtung jeder nimmt seinen eigenen Port

```
## um Grunde macht das auch nur ein port - forward
## Das macht der Trainer nur 1x, dann können alle dort zugreifen
istioctl dashboard grafana --port=3000 --browser=false
```

```
## Jetzt über den Browser öffnen
http://localhost:3000
## Dann Dashboard -> istio -> istio services
```

```
## Lass uns mal Traffic hinschicken vom Client aus
## ip vom ingressgateway from loadBalancer
while true; do curl http://jochen.istio.t3isp.de/api/catalog; sleep .5; done

## Und das das Dashboard nochmal refreshend
##-> General ausklappen
```

## Installation

### Kubernetes mit der Cluster API aufsetzen

#### Komponenten

1. 1x Management-Cluster
2. beliebig viele Workload Cluster, die vom Management Cluster verwaltet werden

#### Voraussetzungen (für Management - Cluster)

- Cluster erstellt mit kubeadm (für Management - Cluster)
- Zugriff zu Cluster auf ./kube/config eingerichtet.
- clusterctl installieren
- kubectl installieren

### Phase 1: Kubernetes-Cluster erstellen und zu Management-Cluster upgraden

#### Schritt 1.1: Cluster erstellen mit kubeadm

- Wir verwenden dafür ein paar selbsterstellte Scripte

```
## Ein Cluster erstellen
cd multi-kubeadmin
./create-multi.sh 1
## Alternativ: Gleich mehrere Cluster für das Training erstellen
## ./create-multi.sh 2
```

#### Schritt 1.2: kubectl runterladen und installieren

```
## Variante 1:
## Systemd needs to work with wsl
## https://devblogs.microsoft.com/commandline/systemd-support-is-now-available-in-wsl/
sudo su -
snap install --classic kubectl
```

```
## Variante 2: Download and install binary (untested)
sudo su -
## https://kubernetes.io/de/docs/tasks/tools/install-kubectl/#installation-der-kubectl-anwendung-mit-curl
curl -LO https://dl.k8s.io/release/${curl -LS https://dl.k8s.io/release/stable.txt}/bin/darwin/amd64/kubectl
chmod +x ./kubectl
mv kubectl /usr/local/bin
```

#### Schritt 1.3: kubeconfig einrichten und Verbindung prüfen

```
## aus dem Onlinesystem unter /etc/kubernetes/admin.conf den Inhalt kopieren
## nach lokal (unprivilegierter Nutzer)
## in
cd
cd .kube
vi config
```

```
kubectl cluster-info
```

#### Schritt 1.4: clusterctl installieren

```
sudo su -
cd /usr/src
curl -L https://github.com/kubernetes-sigs/cluster-api/releases/download/v1.4.2/clusterctl-linux-amd64 -o clusterctl
sudo install -o root -g root -m 0755 clusterctl /usr/local/bin/clusterctl
clusterctl version
```

#### Schritt 1.5: Clusterctl initialisieren (mit dem richtigen provider)

```
## Feature Gate, das für die cluster-api gebraucht wird
export CLUSTER_TOPOLOGY=true
export DIGITALOCEAN_ACCESS_TOKEN=<your-access-token>
export DO_B64ENCODED_CREDENTIALS="$(echo -n "${DIGITALOCEAN_ACCESS_TOKEN}" | base64 | tr -d '\n')"
```

```
## Initialize the management cluster
clusterctl init --infrastructure digitalocean
```

#### Phase 2: Spezielles Kubernetes Images verwenden

- Direkt mit allen Komponenten im Bauch in einer speziellen Version (z.B. 1.28.9)
- die für Kubernetes gebraucht werden

#### Schritt 2.1: install doctl (used to interact with digitalocean)

```
cd
wget https://github.com/digitalocean/doctl/releases/download/v1.107.0/doctl-1.107.0-linux-amd64.tar.gz
tar xf ~/doctl-1.107.0-linux-amd64.tar.gz
sudo mv ~/doctl /usr/local/bin
sudo chmod +x /usr/local/bin/doctl
```

#### Schritt 2.2: install image builder for creating image for digitalocean

```
## if not already done before
export DIGITALOCEAN_ACCESS_TOKEN=<your-access-token>
## as unprivileged user
git clone https://github.com/kubernetes-sigs/image-builder.git
cd image-builder/images/capi
## install dependencies
make deps-do
## show possible builds
make help
## Size of machine will always be 1gb and 1vcpu created in NYC1
make build-do-ubuntu-2404
```

- ACHTUNG: Das dauert eine ganze Weile das bauen

```

digitalocean.ubuntu-2404:
digitalocean.ubuntu-2404: TASK [node : Ensure br_netfilter module is present] *****
digitalocean.ubuntu-2404: changed: [default]
digitalocean.ubuntu-2404:
digitalocean.ubuntu-2404: TASK [node : Persist required kernel modules] *****
digitalocean.ubuntu-2404: changed: [default]
digitalocean.ubuntu-2404:
digitalocean.ubuntu-2404: TASK [node : Set and persist kernel params] *****
digitalocean.ubuntu-2404: changed: [default] => (item={'param': 'net.bridge.bridge-nf-call-iptables', 'val': 1})
digitalocean.ubuntu-2404: changed: [default] => (item={'param': 'net.bridge.bridge-nf-call-ip6tables', 'val': 1})
digitalocean.ubuntu-2404: changed: [default] => (item={'param': 'net.ipv4.ip_forward', 'val': 1})
digitalocean.ubuntu-2404: changed: [default] => (item={'param': 'net.ipv6.conf.all.forwarding', 'val': 1})
digitalocean.ubuntu-2404: changed: [default] => (item={'param': 'net.ipv6.conf.all.disable_ipv6', 'val': 0})
digitalocean.ubuntu-2404: changed: [default] => (item={'param': 'net.ipv4.tcp_congestion_control', 'val': 'bbr'})
digitalocean.ubuntu-2404: changed: [default] => (item={'param': 'vm.overcommit_memory', 'val': 1})
digitalocean.ubuntu-2404: changed: [default] => (item={'param': 'kernel.panic', 'val': 10})
digitalocean.ubuntu-2404: changed: [default] => (item={'param': 'kernel.panic_on_oops', 'val': 1})
digitalocean.ubuntu-2404: changed: [default] => (item={'param': 'fs.inotify.max_user_instances', 'val': 8192})
digitalocean.ubuntu-2404: changed: [default] => (item={'param': 'fs.inotify.max_user_watches', 'val': 524288})
digitalocean.ubuntu-2404:
digitalocean.ubuntu-2404: TASK [node : Ensure auditd is running and comes on at reboot] *****
digitalocean.ubuntu-2404: ok: [default]
digitalocean.ubuntu-2404:
digitalocean.ubuntu-2404: TASK [node : Ensure reverse packet filtering is set as strict] *****
digitalocean.ubuntu-2404: changed: [default]
digitalocean.ubuntu-2404:
digitalocean.ubuntu-2404: TASK [node : Copy udev etcd network tuning rules] *****
digitalocean.ubuntu-2404: changed: [default]

```

---

```

digitalocean.ubuntu-2404: changed: [default]
digitalocean.ubuntu-2404:
digitalocean.ubuntu-2404: TASK [containerd : Delete /opt/cni directory] *****
digitalocean.ubuntu-2404: changed: [default]
digitalocean.ubuntu-2404:
digitalocean.ubuntu-2404: TASK [containerd : Delete /etc/cni directory] *****
digitalocean.ubuntu-2404: changed: [default]
digitalocean.ubuntu-2404:
digitalocean.ubuntu-2404: TASK [containerd : Create unit file directory] *****
digitalocean.ubuntu-2404: changed: [default]
digitalocean.ubuntu-2404:
digitalocean.ubuntu-2404: TASK [containerd : Create containerd memory pressure drop-in file] *****
digitalocean.ubuntu-2404: changed: [default]
digitalocean.ubuntu-2404:
digitalocean.ubuntu-2404: TASK [containerd : Create containerd max tasks drop-in file] *****
digitalocean.ubuntu-2404: changed: [default]
digitalocean.ubuntu-2404:
digitalocean.ubuntu-2404: TASK [containerd : Create containerd http proxy conf file if needed] *****
digitalocean.ubuntu-2404: changed: [default]
digitalocean.ubuntu-2404:
digitalocean.ubuntu-2404: TASK [containerd : Creates containerd config directory] *****
digitalocean.ubuntu-2404: changed: [default]
digitalocean.ubuntu-2404:
digitalocean.ubuntu-2404: TASK [containerd : Copy in containerd config file etc/containerd/config.toml] ***
digitalocean.ubuntu-2404: changed: [default]
digitalocean.ubuntu-2404:
digitalocean.ubuntu-2404: TASK [containerd : Copy in crictl config] *****
digitalocean.ubuntu-2404: changed: [default]

```

### Schritt 2.3: Which kubernetes cluster version is it ?

```

## you need to use exactly the same version for creating your workload cluster
Creating snapshot: Cluster API Kubernetes v1.28.9 on Ubuntu 24.04

```

### Schritt 2.4: Allow Image to be use in Frankfurt Datacenter (FRA1)

```

-> Add to Region FRA1 -> under Manage -> Backups&Snapshots -> Snapshots
Please do this through the web-interface of DigitalOcean
## IF YOU DO NOT DO THIS... Droplets cannot be created because they are in NYC1

```

## Phase 3: Workload - Cluster mit Cluster - API erstellen

### Schritt 3.1 Umgebung zum Ausrollen von ControlNode und Worker vorbereiten

```

## control the datacenter - default nyc1
export DO_REGION=fra1
## control size of machines
## default 1vcpu-1gb
export DO_CONTROL_PLANE_MACHINE_TYPE=s-2vcpu-2gb
export DO_NODE_MACHINE_TYPE=s-2vcpu-2gb
## needed to set up the api provider
export DO_B64ENCODED_CREDENTIALS=""$ \

```

```
echo -n "$DIGITALOCEAN_ACCESS_TOKEN" \  
| base64 \  
| tr -d '\n')
```

### Schritt 3.2 Snapshot-id ausfindig machen

```
doctl compute image list-user
```

```
158401784    Cluster API Kubernetes v1.28.9 on Ubuntu 24.04
```

### Schritt 3.3 Use this snapshot for creation of workload-cluster

```
export DO_CONTROL_PLANE_MACHINE_IMAGE=158401784  
export DO_NODE_MACHINE_IMAGE=158401784
```

### Schritt 3.4 Wir brauchen ein ssh-key

```
## das sollte ein Schlüssel sein, für den wir bereits einen privaten Schlüssel haben  
## und den öffentlichen bei digitalocean hochgeladen haben.  
## Dieser wird dann für die Maschinen verwendet, die hochgezogen werden  
doctl compute ssh-key list
```

```
## wir nehmen den kubernetes key  
42134500    key_training_kubernetes
```

```
## So übergeben wir diesen für das doctl - Tool  
export DO_SSH_KEY_FINGERPRINT=42134500
```

### Schritt 3.5 Cluster.yaml (config) für cluster-api erstellen

```
## Achtung, es muss die gleiche version verwendet werden für die kubernetes version, wie im image  
## das mit dem Image-Builder erstellt wurde
```

```
## Check the variables  
## Show use the necessary env-variables.  
clusterctl generate cluster cluster1 \  
  --infrastructure digitalocean \  
  --target-namespace infra \  
  --kubernetes-version v1.28.9 \  
  --control-plane-machine-count 1 \  
  --worker-machine-count 3 \  
  --list-variables
```

```
## Now create the cluster.yaml file (config to create it)  
## Kuberentes must be the same version as you created the snapshots for do  
## to be used for digitalocean -> creating a cluster there  
clusterctl generate cluster cluster1 \  
  --infrastructure digitalocean \  
  --target-namespace infra \  
  --kubernetes-version v1.28.9 \  
  --control-plane-machine-count 1 \  
  --worker-machine-count 3 \  
  | tee cluster.yaml
```

```
## Create namespace and management cluster for that  
kubect1 create namespace infra
```

```
## and create it  
kubect1 apply --filename cluster.yaml
```

### Schritt 3.5: Wait till controlplane is ready

```
## Achtung das dauert wieder eine ganze Weile,  
## Man kann das im Backend von Digitalocean beobachten  
kubect1 -n infra get kubeadmcontrolplane  
kubect1 -n infra get domachines
```

### Schritt 3.6: Get kubeconfig

```
## When initialized get kubeconfig  
clusterctl --namespace infra \  
  get kubeconfig cluster1 \  
  | tee kubeconfig.yaml
```

### Schritt 3.7. Access new cluster with kubeconfi

```
kubectl --kubeconfig kubeconfig.yaml get ns
## Are all pods ready / of coredns not ;o)
kubectl --kubeconfig kubeconfig.yaml -n kube-system get pods
kubectl --kubeconfig kubeconfig.yaml get nodes
## | | nodes are not ready, because the is no cni-provider installed yet
## v v
```

```
tlnl@pp:~$ kubectl --kubeconfig kubeconfig.yaml get nodes
NAME                                STATUS    ROLES    AGE     VERSION
cluster1-control-plane-wl4j5       NotReady control-plane 8m40s   v1.28.9
cluster1-md-0-45lfj-lbfx           NotReady <none>      3m39s   v1.28.9
cluster1-md-0-45lfj-tkzmf          NotReady <none>      4m4s    v1.28.9
cluster1-md-0-45lfj-zfzh9          NotReady <none>      5m2s    v1.28.9
tlnl@pp:~$ kubectl get -n infra kubeadmcontrolplane
```

### Schritt 3.8. Now install cni -> calico (we will use the operator)

```
kubectl --kubeconfig kubeconfig.yaml create -f https://raw.githubusercontent.com/projectcalico/calico/v3.28.0/manifests/tigera-operator.yaml
## We also want the crd's
kubectl --kubeconfig kubeconfig.yaml create -f https://raw.githubusercontent.com/projectcalico/calico/v3.28.0/manifests/custom-resources.yaml
```

### Schritt 3.9. See the rollout and findout, if everything works

```
## Now watch if everything works smoothly
## everything should be ready
watch kubectl --kubeconfig kubeconfig.yaml get pods -n calico-system
```

```
## + the kube-system coredns pod should also be ready now
kubectl
```

### Schritt 3.10 Installing the CCM (Cloud Controller Manager for digitalocean)

- Important: Before that, the calico-controller will not run, because it cannot get scheduled
- There is a taint on the nodes
  - node.cloudprovider.kubernetes.io/uninitialized
  - There is another taint as well
- These 2 taints will first get removed, once the CCM is installed

```
kubectl --kubeconfig=kubeconfig.yaml apply -f https://raw.githubusercontent.com/digitalocean/digitalocean-cloud-controller-manager/v0.1.54/releases/digitalocean-cloud-controller-manager/v0.1.54.yml
```

### Phase 4: Cleanup

```
## Delete cluster
kubectl -n infra delete cluster cluster1

## Delete management cluster
## after that it probably just will be a normal cluster
kubectl delete cluster
```

### References:

- <https://cluster-api.sigs.k8s.io/user/quick-start.html#install-clusterctl>
- <https://cluster-api.sigs.k8s.io/user/quick-start>
- <https://github.com/kubernetes-sigs/cluster-api-provider-digitalocean/blob/main/docs/getting-started.md>

### Kubernetes mit kubadm aufsetzen (calico)

#### Version

- Ubuntu 20.04 LTS

#### Done for you

- Servers are setup:
  - ssh-running
  - kubeadm, kubelet, kubectl installed
  - containerd - runtime installed
- Installed on all nodes (with cloud-init)

```

#!/bin/bash

groupadd sshadmin
USERS="mysupersecretuser"
SUDO_USER="mysupersecretuser"
PASS="yoursupersecretpass"
for USER in $USERS
do
    echo "Adding user $USER"
    useradd -s /bin/bash --create-home $USER
    usermod -aG sshadmin $USER
    echo "$USER:$PASS" | chpasswd
done

## We can sudo with $SUDO_USER
usermod -aG sudo $SUDO_USER

## 20.04 and 22.04 this will be in the subfolder
if [ -f /etc/ssh/sshd_config.d/50-cloud-init.conf ]
then
    sed -i "s/PasswordAuthentication no/PasswordAuthentication yes/g" /etc/ssh/sshd_config.d/50-cloud-init.conf
fi

### both is needed
sed -i "s/PasswordAuthentication no/PasswordAuthentication yes/g" /etc/ssh/sshd_config

usermod -aG sshadmin root

## TBD - Delete AllowUsers Entries with sed
## otherwise we cannot login by group

echo "AllowGroups sshadmin" >> /etc/ssh/sshd_config
systemctl reload sshd

## Now let us do some generic setup
echo "Installing kubeadm kubelet kubect1"

#### A lot of stuff needs to be done here
#### https://www.linuxtechi.com/install-kubernetes-on-ubuntu-22-04/

## 1. no swap please
swapoff -a
sudo sed -i 's/^(\.*)$/#\1/g' /etc/fstab

## 2. Loading necessary modules
echo "overlay" >> /etc/modules-load.d/containerd.conf
echo "br_netfilter" >> /etc/modules-load.d/containerd.conf
modprobe overlay
modprobe br_netfilter

## 3. necessary kernel settings
echo "net.ipv4.ip_forward = 1" >> /etc/sysctl.d/kubernetes.conf
sysctl --system

## 4. Update the meta-information
apt-get -y update

## 5. Installing container runtime
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o /etc/apt/trusted.gpg.d/docker.add-apt-repository
"deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable" apt-get install -y containerd.io

## 6. Configure containerd
containerd config default > /etc/containerd/config.toml
sed -i 's/SystemdCgroup \= false/SystemdCgroup \= true/g' /etc/containerd/config.toml
systemctl restart containerd
systemctl enable containerd

## 7. Add Kubernetes Repository for Kubernetes
mkdir -m 755 /etc/apt/keyrings
apt-get install -y apt-transport-https ca-certificates curl gpg
curl -fsSL https://pkgs.k8s.io/core:/stable:/$K8S_VERSION/deb/Release.key | gpg --dearmor -o /etc/apt/keyrings/echo "deb [signed-
by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/$K8S_VERSI
# 8. Install kubect1 kubeadm kubect1
apt-get -y update
apt-get install -y kubelet kubeadm kubect1
apt-mark hold -y kubelet kubeadm kubect1

## 9. Install helm

```

```
snap install helm --classic

## Installing nfs-common
apt-get -y install nfs-common
```

### Prerequisites

- 4 Servers setup and reachable through ssh.
- user: 11trainingdo
- pass: PLEASE ask your instructor

```
## Important - Servers are not reachable through
## Domain !! Only IP.
controlplane.tln<nr>.t3isp.de
worker1.tln<nr>.do.t3isp.de
worker2.tln<nr>.do.t3isp.de
worker3.tln<nr>.do.t3isp.de
```

### Step 1: Setup controlnode (login through ssh)

```
## This CIDR is the recommendation for calico
## Other CNI's might be different
CLUSTER_CIDR="192.168.0.0/16"

kubeadm init --pod-network-cidr=$CLUSTER_CIDR && \
  mkdir -p /root/.kube && \
  cp -i /etc/kubernetes/admin.conf /root/.kube/config && \
  chown $(id -u):$(id -g) /root/.kube/config && \
  cp -i /root/.kube/config /tmp/config.kubeadm && \
  chmod o+r /tmp/config.kubeadm

## Copy output of join (needed for workers)
## e.g.
kubeadm join 159.89.99.35:6443 --token rpylp0.rdphpzbavdyx3llz \
  --discovery-token-ca-cert-hash sha256:05d42f2c051a974a27577270e09c77602eeec85523b1815378b815b64cb99932
```

### Step 2: Setup worker1 - node (login through ssh)

```
## use join command from Step 1:
kubeadm join 159.89.99.35:6443 --token rpylp0.rdphpzbavdyx3llz \
  --discovery-token-ca-cert-hash sha256:05d42f2c051a974a27577270e09c77602eeec85523b1815378b815b64cb99932
```

### Step 3: Setup worker2 - node (login through ssh)

```
## use join command from Step 1:
kubeadm join 159.89.99.35:6443 --token rpylp0.rdphpzbavdyx3llz \
  --discovery-token-ca-cert-hash sha256:05d42f2c051a974a27577270e09c77602eeec85523b1815378b815b64cb99932
```

### Step 4: Setup worker3 - node (login through ssh)

```
## use join command from Step 1:
kubeadm join 159.89.99.35:6443 --token rpylp0.rdphpzbavdyx3llz \
  --discovery-token-ca-cert-hash sha256:05d42f2c051a974a27577270e09c77602eeec85523b1815378b815b64cb99932
```

### Step 5: CNI-Setup (calico) on controlnode (login through ssh)

```
kubect1 get nodes

## Output
root@controlplane:~# kubect1 get nodes
NAME           STATUS    ROLES          AGE      VERSION
controlplane   NotReady  control-plane  6m27s    v1.28.6
worker1        NotReady  <none>         3m18s    v1.28.6
worker2        NotReady  <none>         2m10s    v1.28.6
worker3        NotReady  <none>         60s      v1.28.6

## Installing calico CNI
kubect1 create -f https://raw.githubusercontent.com/projectcalico/calico/v3.27.0/manifests/tigera-operator.yaml
kubect1 create -f https://raw.githubusercontent.com/projectcalico/calico/v3.27.0/manifests/custom-resources.yaml
kubect1 get ns
kubect1 -n calico-system get all
kubect1 -n calico-system get pods -o wide -w

## After if all pods are up and running -> CTRL + C
```



```
kubectl -n calico-system get pods -o wide
## all nodes should be ready now
kubectl get nodes -o wide
```

```
## Output
root@controlplane:~# kubectl get nodes
NAME           STATUS    ROLES    AGE   VERSION
controlplane   Ready     control-plane   14m   v1.28.6
worker1        Ready     <none>        11m   v1.28.6
worker2        Ready     <none>        10m   v1.28.6
worker3        Ready     <none>        9m9s   v1.28.6
```

## Kubernetes - Misc

### Wann wird podIP vergeben ?

#### Example (that does work)

```
## Show the pods that are running
kubectl get pods

## Synopsis (most simplistic example)
## kubectl run NAME --image=IMAGE_EG_FROM_DOCKER
## example
kubectl run nginx --image=nginx:1.23

kubectl get pods
## on which node does it run ?
kubectl get pods -o wide
```

#### Example (that does not work)

```
kubectl run foo2 --image=foo2
## ImageErrPull - Image konnte nicht geladen werden
kubectl get pods
## Weitere status - info
kubectl describe pods foo2
```

#### Ref:

- <https://kubernetes.io/docs/reference/generated/kubectl/kubectl-commands#run>

### Bash completion installieren

#### Walkthrough

```
## Eventuell, wenn bash-completion nicht installiert ist.
apt install bash-completion
source /usr/share/bash-completion/bash_completion
## is it installed properly
type _init_completion

## activate for all users
kubectl completion bash | sudo tee /etc/bash_completion.d/kubectl > /dev/null

## verifizieren - neue login shell
su -

## zum Testen
kubectl g<TAB>
kubectl get
```

#### Alternative für k als alias für kubectl

```
source <(kubectl completion bash)
complete -F __start_kubectl k
```

#### Reference

- <https://kubernetes.io/docs/tasks/tools/included/optional-kubectl-configs-bash-linux/>

### Remote-Verbindung zu Kubernetes (microk8s) einrichten

#### Step 1: Install kubectl on local machine (or jump-server)

```
## on CLIENT install kubectl
sudo snap install kubectl --classic
```

## Step 2: configure kubectl

```
## On MASTER -server get config
## als root
microk8s config > /tmp/config
cat /tmp/config
```

```
## Optional or simply copy & paste
## Download (scp config file) and store in .kube - folder
```

```
cd
mkdir .kube
cd .kube # Wichtig: config muss nachher im verzeichnis .kube liegen
## scp kurs@master_server:/path/to/remote_config config
## z.B.
scp kurs@192.168.56.102:/home/kurs/remote_config config
## oder benutzer 11trainingdo
scp 11trainingdo@192.168.56.102:/home/11trainingdo/remote_config config

#### Evtl. IP-Adresse in config zum Server aendern

## Ultimate 1. Test auf CLIENT
kubectl cluster-info

## or if using kubectl or alias
kubectl get pods

## if you want to use a different kube config file, you can do like so
kubectl --kubeconfig /home/myuser/.kube/myconfig
```

## vim support for yaml

### Ubuntu (im Unterverzeichnis /etc/vim/vimrc.local - systemweit)

```
hi CursorColumn cterm=NONE ctermbg=lightred ctermfg=white
autocmd FileType y?ml setlocal ts=2 sts=2 sw=2 ai number expandtab cursorline cursorcolumn
```

## Testen

```
vim test.yml
Eigenschaft: <return> # springt eingerückt in die nächste Zeile um 2 spaces eingerückt


## evtl funktioniert vi test.yml auf manchen Systemen nicht, weil kein vim (vi improved)
```

## Kubernetes - Netzwerk (CNI's) / Mesh

### Netzwerk Interna

### Network Namespace for each pod

### Overview



#### General

- Each pod will have its own network namespace
  - with routing, network devices
- Connection to default namespace to host is done through veth - Link to bridge on host network
  - similar like on docker to docker0

Each container is connected to the bridge via a veth-pair. This interface pair functions like a virtual point-to-point ethernet connection and connects the network namespaces of the containers with the network namespace of the host

- Every container is in the same Network Namespace, so they can communicate through localhost
  - Example with hashicorp/http-echo container 1 and busybox container 2 ?

#### Pod-To-Pod Communication (across nodes)


##### Prerequisites

- pods on a single node as well as pods on a topological remote can establish communication at all times
- Each pod receives a unique IP address, valid anywhere in the cluster. Kubernetes requires this address to not be subject to network address translation (NAT)
- Pods on the same node through virtual bridge (see image above)

##### General (what needs to be done) - and could be done manually

- local bridge networks of all nodes need to be connected
- there needs to be an IPAM (IP-Address Management) so addresses are only used once
- The need to be routes so, that each bridge can communicate with the bridge on the other network
- Plus: There needs to be a rule for incoming network
- Also: A tunnel needs to be set up to the outside world.

##### General - Pod-to-Pod Communication (across nodes) - what would need to be done



#### General - Pod-to-Pod Communication (side-note)

- This could of course be done manually, but it is too complex
- So Kubernetes has created an Interface, which is well defined
  - The interface is called CNI (common network interface)
  - Functionally is achieved through Network Plugin (which use this interface)
    - e.g. calico / cilium / weave net / flannel

#### CNI

- CNI only handles network connectivity of container and the cleanup of allocated resources (i.e. IP addresses) after containers have been deleted (garbage collection) and therefore is lightweight and quite easy to implement.
- There are some basic libraries within CNI which do some basic stuff.

#### Hidden Pause Container

##### What is for ?

- Holds the network - namespace for the pod
- Gets started first and falls asleep later
- Will still be there, when the other containers die

```
cd
mkdir -p manifests
cd manifests
mkdir pausetest
cd pausetest
nano 01-nginx.yml

## vi nginx-static.yml

apiVersion: v1
kind: Pod
metadata:
  name: nginx-pausetest
  labels:
    webserver: nginx:1.21
spec:
  containers:
    - name: web
      image: nginx

kubectl apply -f .

ctr -n k8s.io c list | grep pause
```

#### References

- <https://www.inovex.de/de/blog/kubernetes-networking-part-1-en/>
- <https://www.inovex.de/de/blog/kubernetes-networking-2-calico-cilium-weavenet/>

## Übersicht Netzwerke

### CNI

- Common Network Interface
- Feste Definition, wie Container mit Netzwerk-Bibliotheken kommunizieren

### Docker - Container oder andere

- Container wird hochgefahren -> über CNI -> zieht Netzwerk - IP hoch.
- Container wird runtergefahren -> über CNI -> Netzwerk - IP wird released

### Welche gibt es ?

- Flannel
- Canal
- Calico
- Cilium
- Weave Net

### Flannel

#### Overlay - Netzwerk

- virtuelles Netzwerk was sich oben drüber und eigentlich auf Netzwerkebene nicht existiert
- VXLAN

#### Vorteile

- Guter einfacher Einstieg
- reduziert auf eine Binary flanneld

#### Nachteile

- keine Firewall - Policies möglich
- keine klassischen Netzwerk-Tools zum Debuggen möglich.

### Canal

#### General

- Auch ein Overlay - Netzwerk
- Unterstützt auch policies

### Calico

#### Generell

- klassische Netzwerk (BGP)

#### Vorteile gegenüber Flannel

- Policy über Kubernetes Object (NetworkPolicies)

#### Vorteile

- ISTIO integrierbar (Mesh - Netz)
- Performance etwas besser als Flannel (weil keine Encapsulation)

#### Referenz

- <https://projectcalico.docs.tigera.io/security/calico-network-policy>

### Cilium

### Weave Net

- Ähnlich calico
- Verwendet overlay netzwerk
- Sehr stabil bzgl IPV4/IPV6 (Dual Stack)
- Sehr grosses Feature-Set
- mit das älteste Plugin

### microk8s Vergleich

- <https://microk8s.io/compare>

```

snap.microk8s.daemon-flanneld
Flannel is a CNI which gives a subnet to each host for use with container runtimes.

Flanneld runs if ha-cluster is not enabled. If ha-cluster is enabled, calico is run instead.

The flannel daemon is started using the arguments in ${SNAP_DATA}/args/flanneld. For more information on the configuration, see
the flannel documentation.

```

### IPV4/IPV6 Dualstack

- <https://kubernetes.io/docs/concepts/services-networking/dual-stack/>

## Ingress controller in microk8s aktivieren

### Aktivieren

```
microk8s enable ingress
```

### Referenz

- <https://microk8s.io/docs/addon-ingress>

## Kubernetes - Ingress

### ingress mit ssl absichern

## Kubernetes - Wartung / Debugging

### kubectl drain/uncordon

```
## Achtung, bitte keine pods verwenden, dies können "ge"-drained (ausgetrocknet) werden
kubectl drain <node-name>
z.B.
## Daemonsets ignorieren, da diese nicht gelöscht werden
kubectl drain n17 --ignore-daemonsets

## Alle pods von replicaset werden jetzt auf andere nodes verschoben
## Ich kann jetzt wartungsarbeiten durchführen

## Wenn fertig bin:
kubectl uncordon n17

## Achtung: deployments werden nicht neu ausgerollt, dass muss ich anstossen.
## z.B.
kubectl rollout restart deploy/webserver
```

### Alte manifeste konvertieren mit convert plugin

#### What is about?

- Plugins needs to be installed seperately on Client (or where you have your manifests)

#### Walkthrough

```
curl -LO "https://dl.k8s.io/release/$(curl -L -s https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl-convert"
## Validate the checksum
curl -LO "https://dl.k8s.io/$(curl -L -s https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl-convert.sha256"
echo "${kubectl-convert.sha256} kubectl-convert" | sha256sum --check
## install
sudo install -o root -g root -m 0755 kubectl-convert /usr/local/bin/kubectl-convert

## Does it work
kubectl convert --help

## Works like so
## Convert to the newest version
## kubectl convert -f pod.yaml
```

### Reference

- <https://kubernetes.io/docs/tasks/tools/install-kubectl-linux/#install-kubectl-convert-plugin>

### Curl from pod api-server

<https://nieldw.medium.com/curling-the-kubernetes-api-server-d7675cfc398c>

## Kubernetes Praxis API-Objekte

### kubectl example with run

#### Example (that does work)

```
## Show the pods that are running
kubectl get pods

## Synopsis (most simplistic example
## kubectl run NAME --image=IMAGE_EG_FROM_DOCKER
## example
kubectl run nginx --image=nginx:1.23
```

```
kubectl get pods
## on which node does it run ?
kubectl get pods -o wide
```

#### Example (that does not work)

```
kubectl run foo2 --image=foo2
## ImageErrPull - Image konnte nicht geladen werden
kubectl get pods
## Weitere status - info
kubectl describe pods foo2
```

#### Ref:

- <https://kubernetes.io/docs/reference/generated/kubect/kubectl-commands#run>

#### Ingress Controller auf Digitalocean (doks) mit helm installieren

##### Basics

- works for all plattform with helm if no ingresscontroller ist present
- if you have ingress - objekts and no ingresscontroller nothing works

##### Prerequisites

- kubectl must be set up

##### Walkthrough (Setup Ingress Controller)

```
helm repo add ingress-nginx https://kubernetes.github.io/ingress-nginx
helm repo update

## helm show values ingress-nginx/ingress-nginx

## It will be setup with type loadbalancer - so waiting to retrieve an ip from the external loadbalancer
## This will take a little.
helm install nginx-ingress ingress-nginx/ingress-nginx --namespace ingress --create-namespace --set
controller.publishService.enabled=true

## See when the external ip comes available
kubectl -n ingress get all
kubectl --namespace ingress get services -o wide -w nginx-ingress-ingress-nginx-controller

## Output
NAME                                TYPE           CLUSTER-IP    EXTERNAL-IP    PORT(S)                                AGE
SELECTOR
nginx-ingress-ingress-nginx-controller LoadBalancer  10.245.78.34   157.245.20.222 80:31588/TCP,443:30704/TCP 4m39s
app.kubernetes.io/component=controller,app.kubernetes.io/instance=nginx-ingress,app.kubernetes.io/name=ingress-nginx

## Now setup wildcard - domain for training purpose
## inwx.com
*.lab1.t3isp.de A 157.245.20.222
```

#### Documentation for default ingress nginx

- <https://kubernetes.github.io/ingress-nginx/user-guide/nginx-configuration/configmap/>

#### Beispiel Ingress

##### Prerequisites

```
## Ingress Controller muss aktiviert sein
microk8s enable ingress
```

##### Walkthrough

##### Schritt 1:

```
cd
mkdir -p manifests
cd manifests
mkdir abi
cd abi
```

```
## apple.yml
## vi apple.yml
kind: Pod
apiVersion: v1
metadata:
```

```

name: apple-app
labels:
  app: apple
spec:
  containers:
    - name: apple-app
      image: hashicorp/http-echo
      args:
        - "-text=apple"
---

kind: Service
apiVersion: v1
metadata:
  name: apple-service
spec:
  selector:
    app: apple
  ports:
    - protocol: TCP
      port: 80
      targetPort: 5678 # Default port for image

```

```
kubectl apply -f apple.yml
```

```

## banana
## vi banana.yml
kind: Pod
apiVersion: v1
metadata:
  name: banana-app
  labels:
    app: banana
spec:
  containers:
    - name: banana-app
      image: hashicorp/http-echo
      args:
        - "-text=banana"
---

kind: Service
apiVersion: v1
metadata:
  name: banana-service
spec:
  selector:
    app: banana
  ports:
    - port: 80
      targetPort: 5678 # Default port for image

```

```
kubectl apply -f banana.yml
```

## Schritt 2:

```

## Ingress
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: example-ingress
  annotations:
    ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx
  rules:
    - http:
        paths:
          - path: /apple
            backend:
              serviceName: apple-service
              servicePort: 80
          - path: /banana
            backend:
              serviceName: banana-service
              servicePort: 80

```



```
## ingress
kubectl apply -f ingress.yml
kubectl get ing
```

## Reference

- <https://matthewpalmer.net/kubernetes-app-developer/articles/kubernetes-ingress-guide-nginx-example.html>

## Find the problem

```
## Hints

## 1. Which resources does our version of kubectl support
## Can we find Ingress as "Kind" here.
kubectl api-resources

## 2. Let's see, how the configuration works
kubectl explain --api-version=networking.k8s.io/v1 ingress.spec.rules.http.paths.backend.service

## now we can adjust our config
```

## Solution

```
## in kubernetes 1.22.2 - ingress.yml needs to be modified like so.
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example-ingress
  annotations:
    ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx
  rules:
  - http:
      paths:
      - path: /apple
        pathType: Prefix
        backend:
          service:
            name: apple-service
            port:
              number: 80
      - path: /banana
        pathType: Prefix
        backend:
          service:
            name: banana-service
            port:
              number: 80
```

## Achtung: Ingress mit Helm - annotations

## Permanente Weiterleitung mit Ingress

## Example

```
## redirect.yml
apiVersion: v1
kind: Namespace
metadata:
  name: my-namespace
---

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  annotations:
    nginx.ingress.kubernetes.io/permanent-redirect: https://www.google.de
    nginx.ingress.kubernetes.io/permanent-redirect-code: "308"
  creationTimestamp: null
  name: destination-home
  namespace: my-namespace
spec:
  rules:
  - host: web.training.local
    http:
```

```
paths:
- backend:
    service:
      name: http-svc
      port:
        number: 80
    path: /source
    pathType: ImplementationSpecific
```

Achtung: host-eintrag auf Rechner machen, von dem aus man zugreift

```
/etc/hosts
45.23.12.12 web.training.local
```

```
curl -I http://web.training.local/source
HTTP/1.1 308
Permanent Redirect
```

### Umbauen zu google ;o)

This annotation allows to return a permanent redirect instead of sending data to the upstream. For example `nginx.ingress.kubernetes.io/permanent-redirect: https://www.google.com` would redirect everything to Google.

### Refs:

- <https://github.com/kubernetes/ingress-nginx/blob/main/docs/user-guide/nginx-configuration/annotations.md#permanent-redirect>
- 

### ConfigMap Example

#### Schritt 1: configmap vorbereiten

```
cd
mkdir -p manifests
cd manifests
mkdir configmaptests
cd configmaptests
nano 01-configmap.yml
```

```
### 01-configmap.yml
kind: ConfigMap
apiVersion: v1
metadata:
  name: example-configmap
data:
  # als Wertepaare
  database: mongodb
  database_uri: mongodb://localhost:27017
```

```
kubectl apply -f 01-configmap.yml
kubectl get cm
kubectl get cm -o yaml
```

#### Schritt 2: Beispiel als Datei

```
nano 02-pod.yml
```

```
kind: Pod
apiVersion: v1
metadata:
  name: pod-mit-configmap

spec:
  # Add the ConfigMap as a volume to the Pod
  volumes:
    # `name` here must match the name
    # specified in the volume mount
    - name: example-configmap-volume
      # Populate the volume with config map data
      configMap:
        # `name` here must match the name
        # specified in the ConfigMap's YAML
        name: example-configmap

  containers:
    - name: container-configmap
```

```

image: nginx:latest
# Mount the volume that contains the configuration data
# into your container filesystem
volumeMounts:
  # `name` here must match the name
  # from the volumes section of this pod
  - name: example-configmap-volume
    mountPath: /etc/config

```

```
kubectl apply -f 02-pod.yml
```

```

##Jetzt schauen wir uns den Container/Pod mal an
kubectl exec pod-mit-configmap -- ls -la /etc/config
kubectl exec -it pod-mit-configmap -- bash
## ls -la /etc/config

```

### Schritt 3: Beispiel. ConfigMap als env-variablen

```
nano 03-pod-mit-env.yml
```

```

## 03-pod-mit-env.yml
kind: Pod
apiVersion: v1
metadata:
  name: pod-env-var
spec:
  containers:
    - name: env-var-configmap
      image: nginx:latest
      envFrom:
        - configMapRef:
            name: example-configmap

```

```
kubectl apply -f 03-pod-mit-env.yml
```

```

## und wir schauen uns das an
##Jetzt schauen wir uns den Container/Pod mal an
kubectl exec pod-env-var -- env
kubectl exec -it pod-env-var -- bash
## env

```

### Reference:

- <https://matthewpalmer.net/kubernetes-app-developer/articles/ultimate-configmap-guide-kubernetes.html>

### Configmap MariaDB my.cnf

#### configmap zu fuss

```
vi mariadb-config2.yml
```

```

kind: ConfigMap
apiVersion: v1
metadata:
  name: example-configmap
data:
  # als Wertepaare
  database: mongodb
  my.cnf: |
[mysqld]
slow_query_log = 1
innodb_buffer_pool_size = 1G

```

```
kubectl apply -f .
```

```

##deploy.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mariadb-deployment
spec:
  selector:
    matchLabels:
      app: mariadb

```

```

replicas: 1
template:
  metadata:
    labels:
      app: mariadb
  spec:
    containers:
      - name: mariadb-cont
        image: mariadb:latest
        envFrom:
          - configMapRef:
              name: mariadb-configmap

        volumeMounts:
          - name: example-configmap-volume
            mountPath: /etc/my

    volumes:
      - name: example-configmap-volume
        configMap:
          name: example-configmap

```

```
kubectl apply -f .
```

## Helm (Kubernetes Paketmanager)

### Helm Grundlagen

#### Wo ?

```
artifacts helm
```

- <https://artifacthub.io/>

### Komponenten

```

Chart - beinhaltet Beschreibung und Komponenten
tar.gz - Format
oder Verzeichnis

```

```

Wenn wir ein Chart ausführen wird eine Release erstellen
(parallel: image -> container, analog: chart -> release)

```

### Installation

```

## Beispiel ubuntu
## snap install --classic helm

## Cluster muss vorhanden, aber nicht notwendig wo helm installiert

## Voraussetzung auf dem Client-Rechner (helm ist nichts als anderes als ein Client-Programm)
Ein lauffähiges kubectl auf dem lokalen System (welches sich mit dem Cluster verbinden kann).
-> saubere -> .kube/config

## Test
kubectl cluster-info

```

### Helm Warum ?

```

Ein Paket für alle Komponenten
Einfaches Installieren, Updaten und deinstallieren
Feststehende Struktur

```

### Helm Example

#### Prerequisites

- kubectl needs to be installed and configured to access cluster
- Good: helm works as unprivileged user as well - Good for our setup
- install helm on ubuntu (client) as root: snap install --classic helm
  - this installs helm3
- Please only use: helm3. No server-side components needed (in cluster)
  - Get away from examples using helm2 (hint: helm init) - uses tiller

#### Simple Walkthrough (Example 0)

```
## Repo hinzufügen
helm repo add bitnami https://charts.bitnami.com/bitnami
## gecachte Informationen aktualisieren
helm repo update

helm search repo bitnami
## helm install release-name bitnami/mysql
helm install my-mysql bitnami/mysql
## Chart runterziehen ohne installieren
## helm pull bitnami/mysql

## Release anzeigen zu lassen
helm list

## Status einer Release / Achtung, heisst nicht unbedingt nicht, dass pod läuft
helm status my-mysql

## weitere release installieren
## helm install neuer-release-name bitnami/mysql
```

### Under the hood

```
## Helm speichert Informationen über die Releases in den Secrets
kubectl get secrets | grep helm
```

### Example 1: - To get know the structure

```
helm repo add bitnami https://charts.bitnami.com/bitnami
helm search repo bitnami
helm repo update
helm pull bitnami/mysql
tar xzvf mysql-9.0.0.tgz
```

### Example 2: We will setup mysql without persistent storage (not helpful in production ;o)

```
helm repo add bitnami https://charts.bitnami.com/bitnami
helm search repo bitnami
helm repo update

helm install my-mysql bitnami/mysql
```

### Example 2 - continue - fehlerbehebung

```
helm uninstall my-mysql
## Install with persistentStorage disabled - Setting a specific value
helm install my-mysql --set primary.persistence.enabled=false bitnami/mysql

## just as notice
## helm uninstall my-mysql
```

### Example 2b: using a values file

```
## mkdir helm-mysql
## cd helm-mysql
## vi values.yml
primary:
  persistence:
    enabled: false
```

```
helm uninstall my-mysql
helm install my-mysql bitnami/mysql -f values.yml
```

### Example 3: Install wordpress

```
helm repo add bitnami https://charts.bitnami.com/bitnami
helm install my-wordpress \
  --set wordpressUsername=admin \
  --set wordpressPassword=password \
  --set mariadb.auth.rootPassword=secretpassword \
  bitnami/wordpress
```

#### Example 4: Install Wordpress with values and auth

```
## mkdir helm-mysql
## cd helm-mysql
## vi values.yml
persistence:
  enabled: false

wordpressUsername: admin
wordpressPassword: password
mariadb:
  primary:
    persistence:
      enabled: false

auth:
  rootPassword: secretpassword
```

```
helm uninstall my-wordpress
helm install my-wordpress bitnami/wordpress -f values
```

#### Referenced

- <https://github.com/bitnami/charts/tree/master/bitnami/mysql/#installing-the-chart>
- <https://helm.sh/docs/intro/quickstart/>

## Kubernetes - RBAC

### Nutzer einrichten microk8s ab kubernetes 1.25

#### Enable RBAC in microk8s

```
## This is important, if not enable every user on the system is allowed to do everything
microk8s enable rbac
```

#### Schritt 1: Nutzer-Account auf Server anlegen und secret anlegen / in Client

```
cd
mkdir -p manifests/rbac
cd manifests/rbac
```

##### Mini-Schritt 1: Definition für Nutzer

```
## vi service-account.yml
apiVersion: v1
kind: ServiceAccount
metadata:
  name: training
  namespace: default
```

```
kubectl apply -f service-account.yml
```

##### Mini-Schritt 1.5: Secret erstellen

- From Kubernetes 1.25 tokens are not created automatically when creating a service account (sa)
- You have to create them manually with annotation attached
- <https://kubernetes.io/docs/reference/access-authn-authz/service-accounts-admin/#create-token>

```
## vi secret.yml
apiVersion: v1
kind: Secret
type: kubernetes.io/service-account-token
metadata:
  name: trainingtoken
  annotations:
    kubernetes.io/service-account.name: training
```

```
kubectl apply -f .
```

##### Mini-Schritt 2: ClusterRole festlegen - Dies gilt für alle namespaces, muss aber noch zugewiesen werden

```
### Bevor sie zugewiesen ist, funktioniert sie nicht - da sie keinem Nutzer zugewiesen ist
```

```
## vi pods-clusterrole.yml
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: pods-clusterrole
rules:
- apiGroups: [""] # "" indicates the core API group
  resources: ["pods"]
  verbs: ["get", "watch", "list", "create"]
```

```
kubectl apply -f pods-clusterrole.yml
```

### Mini-Schritt 3: Die ClusterRolle den entsprechenden Nutzern über RoleBinding zu ordnen

```
## vi rb-training-ns-default-pods.yml
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: rolebinding-ns-default-pods
  namespace: default
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: pods-clusterrole
subjects:
- kind: ServiceAccount
  name: training
  namespace: default
```

```
kubectl apply -f rb-training-ns-default-pods.yml
```

### Mini-Schritt 4: Testen (klappt der Zugang)

```
kubectl auth can-i get pods -n default --as system:serviceaccount:default:training
```

## Schritt 2: Context anlegen / Credentials auslesen und in kubeconfig hinterlegen (ab Kubernetes-Version 1.25.)

### Mini-Schritt 1: kubeconfig setzen

```
kubectl config set-context training-ctx --cluster microk8s-cluster --user training

## extract name of the token from here

TOKEN=`kubectl get secret trainingtoken -o jsonpath='{.data.token}' | base64 --decode`
echo $TOKEN
kubectl config set-credentials training --token=$TOKEN
kubectl config use-context training-ctx

## Hier reichen die Rechte nicht aus
kubectl get deploy
## Error from server (Forbidden): pods is forbidden: User "system:serviceaccount:kube-system:training" cannot list # resource
"pods" in API group "" in the namespace "default"
```

### Mini-Schritt 2:

```
kubectl config use-context training-ctx
kubectl get pods
```

### Mini-Schritt 3: Zurück zum alten Default-Context

```
kubectl config get-contexts
```

CURRENT	NAME	CLUSTER	AUTHINFO	NAMESPACE
	microk8s	microk8s-cluster	admin2	
*	training-ctx	microk8s-cluster	training2	

```
kubectl config use-context microk8s
```

### Refs:

- <https://docs.oracle.com/en-us/iaas/Content/ContEng/Tasks/contengaddingserviceaccttoken.htm>
- <https://microk8s.io/docs/multi-user>
- <https://faun.pub/kubernetes-rbac-use-one-role-in-multiple-namespaces-d1d08bb08286>

## Ref: Create Service Account Token

- <https://kubernetes.io/docs/reference/access-authn-authz/service-accounts-admin/#create-token>

## Tipps&Tricks zu Deployment - Rollout

### Warum

Rückgängig machen von deploys, Deploys neu unstossen.  
(Das sind die wichtigsten Fähigkeiten)

### Beispiele

```
## Deployment nochmal durchführen
## z.B. nach kubectl uncordon n12.training.local
kubectl rollout restart deploy nginx-deployment

## Rollout rückgängig machen
kubectl rollout undo deploy nginx-deployment
```

## Kustomize

### Kustomize Overlay Beispiel

#### Konzept Overlay

- Base + Overlay = Gepatchtes manifest
- Sachen patchen.
- Die werden drübergelegt.

#### Example 1: Walkthrough

```
## Step 1:
## Create the structure
## kustomize-example1
## L base
## | - kustomization.yml
## L overlays
##.   L dev
##     - kustomization.yml
##.   L prod
##     - kustomization.yml
cd; mkdir -p manifests/kustomize-example1/base; mkdir -p manifests/kustomize-example1/overlays/prod; cd manifests/kustomize-example1
```

```
## Step 2: base dir with files
## now create the base kustomization file
## vi base/kustomization.yml
resources:
- service.yml
```

```
## Step 3: Create the service - file
## vi base/service.yml
kind: Service
apiVersion: v1
metadata:
  name: service-app
spec:
  type: ClusterIP
  selector:
    app: simple-app
  ports:
    - name: http
      port: 80
```

```
## See how it looks like
kubectl kustomize ./base
```

```
## Step 4: create the customization file accordingly
##vi overlays/prod/kustomization.yml
bases:
- ../../base
patches:
- service-ports.yaml
```

```
## Step 5: create overlay (patch files)
## vi overlays/prod/service-ports.yaml
```



```
kind: Service
apiVersion: v1
metadata:
  #Name der zu patchenden Ressource
  name: service-app
spec:
  # Changed to Nodeport
  type: NodePort
  ports: #Die Porteinstellungen werden überschrieben
  - name: https
    port: 443
```

```
## Step 6:
kubectl kustomize overlays/prod

## or apply it directly
kubectl apply -k overlays/prod/
```

```
## Step 7:
## mkdir -p overlays/dev
## vi overlays/dev/kustomization
bases:
- ../../base
```

```
## Step 8:
## statt mit der base zu arbeiten
kubectl kustomize overlays/dev
```

#### Example 2: Advanced Patching with patchesJson6902 (You need to have done example 1 firstly)

```
## Schritt 1:
## Replace overlays/prod/kustomization.yml with the following syntax
bases:
- ../../base
patchesJson6902:
- target:
    version: v1
    kind: Service
    name: service-app
  path: service-patch.yaml
```

```
## Schritt 2:
## vi overlays/prod/service-patch.yaml
- op: remove
  path: /spec/ports
  value:
    - name: http
      port: 80
- op: add
  path: /spec/ports
  value:
    - name: https
      port: 443
```

```
## Schritt 3:
kubectl kustomize overlays/prod
```

#### Special Use Case: Change the metadata.name

```
## Same as Example 2, but patch-file is a bit different
## vi overlays/prod/service-patch.yaml
- op: remove
  path: /spec/ports
  value:
    - name: http
      port: 80

- op: add
  path: /spec/ports
  value:
    - name: https
      port: 443

- op: replace
  path: /metadata/name
  value: svc-app-test
```

```
kubectl kustomize overlays/prod
```

#### Ref:

- <https://blog.ordix.de/kubernetes-anwendungen-mit-kustomize>

#### Helm mit kustomize verheiraten

## Kubernetes - Tipps & Tricks

### Kubernetes Debuggen ClusterIP/PodIP

#### Situation

- Kein Zugriff auf die Nodes, zum Testen von Verbindungen zu Pods und Services über die PodIP/ClusterIP

#### Lösung

```
## Wir starten eine Busybox und fragen per wget und port ab
## busytester ist der name
## long version
kubectl run -it --rm --image=busybox busytester
## wget <pod-ip-des-ziels>
## exit

## quick and dirty
kubectl run -it --rm --image=busybox busytester -- wget <pod-ip-des-ziels>
```

### Debugging pods

#### How ?

1. Which pod is in charge
2. Problems when starting: `kubectl describe po mypod`
3. Problems while running: `kubectl logs mypod`

### Taints und Tolerations

#### Taints

```
Taints schliessen auf einer Node alle Pods aus, die nicht bestimmte taints haben:

Möglichkeiten:

o Sie werden nicht gescheduled - NoSchedule
o Sie werden nicht executed - NoExecute
o Sie werden möglichst nicht gescheduled. - PreferNoSchedule
```

#### Tolerations

```
Tolerations werden auf Pod-Ebene vergeben:
tolerations:

Ein Pod kann (wenn es auf einem Node taints gibt), nur
gescheduled bzw. ausgeführt werden, wenn er die
Labels hat, die auch als
Taints auf dem Node vergeben sind.
```

### Walkthrough

#### Step 1: Cordon the other nodes - scheduling will not be possible there

```
## Cordon nodes n11 and n111
## You will see a taint here
kubectl cordon n11
kubectl cordon n111
kubectl describe n111 | grep -i taint
```

#### Step 2: Set taint on first node

```
kubectl taint nodes n1 gpu=true:NoSchedule
```

#### Step 3

```
cd
mkdir -p manifests
```

```
cd manifests
mkdir tainttest
cd tainttest
nano 01-no-tolerations.yml
```

```
##vi 01-no-tolerations.yml
apiVersion: v1
kind: Pod
metadata:
  name: nginx-test-no-tol
  labels:
    env: test-env
spec:
  containers:
    - name: nginx
      image: nginx:1.21
```

```
kubectl apply -f .
kubectl get po nginx-test-no-tol
kubectl get describe nginx-test-no-tol
```

#### Step 4:

```
## vi 02-nginx-test-wrong-tol.yml
apiVersion: v1
kind: Pod
metadata:
  name: nginx-test-wrong-tol
  labels:
    env: test-env
spec:
  containers:
    - name: nginx
      image: nginx:latest
  tolerations:
    - key: "cpu"
      operator: "Equal"
      value: "true"
      effect: "NoSchedule"
```

```
kubectl apply -f .
kubectl get po nginx-test-wrong-tol
kubectl describe po nginx-test-wrong-tol
```

#### Step 5:

```
## vi 03-good-tolerations.yml
apiVersion: v1
kind: Pod
metadata:
  name: nginx-test-good-tol
  labels:
    env: test-env
spec:
  containers:
    - name: nginx
      image: nginx:latest
  tolerations:
    - key: "gpu"
      operator: "Equal"
      value: "true"
      effect: "NoSchedule"
```

```
kubectl apply -f .
kubectl get po nginx-test-good-tol
kubectl describe po nginx-test-good-tol
```

#### Taints rausnehmen

```
kubectl taint nodes n1 gpu:true:NoSchedule-
```

#### uncordon other nodes

```
kubectl uncordon n11
kubectl uncordon n111
```

## References

- [Doku Kubernetes Taints and Tolerations](#)
- <https://blog.kubecost.com/blog/kubernetes-taints/>

## pod aus deployment bei config - Änderung neu ausrollen

- <https://github.com/stakater/Reloader>

## Kubernetes Advanced

### Curl api-server kubernetes aus pod heraus

<https://nielddw.medium.com/curling-the-kubernetes-api-server-d7675cfc398c>

## Kubernetes - Documentation

### Documentation zu microk8s plugins/addons

- <https://microk8s.io/docs/addons>

### Shared Volumes - Welche gibt es ?

- <https://kubernetes.io/docs/concepts/storage/volumes/>

## Kubernetes - Hardening

### Kubernetes Tipps Hardening

### PSA (Pod Security Admission)

```
Policies defined by namespace.  
e.g. not allowed to run container as root.  
  
Will complain/deny when creating such a pod with that container type
```

### Möglichkeiten in Pods und Containern

```
## für die Pods  
kubectl explain pod.spec.securityContext  
kubectl explain pod.spec.containers.securityContext
```

### Example (seccomp / security context)

```
A. seccomp - profile  
https://github.com/docker/docker/blob/master/profiles/seccomp/default.json
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: audit-pod  
  labels:  
    app: audit-pod  
spec:  
  securityContext:  
    seccompProfile:  
      type: Localhost  
      localhostProfile: profiles/audit.json  
  
  containers:  
  
  - name: test-container  
    image: hashicorp/http-echo:0.2.3  
    args:  
      - "-text=just made some syscalls!"  
    securityContext:  
      allowPrivilegeEscalation: false
```

### SecurityContext (auf Pod Ebene)

```
kubectl explain pod.spec.containers.securityContext
```

### NetworkPolicy

```
## Firewall Kubernetes
```

### Kubernetes Security Admission Controller Example

### Seit: 1.2.22 Pod Security Admission

- 1.2.22 - Alpha - D.h. ist noch nicht aktiviert und muss als Feature Gate aktiviert (Kind)
- 1.2.23 - Beta -> d.h. evtl. aktiviert

### Vorgefertigte Regelwerke

- privileges - keinerlei Einschränkungen
- baseline - einige Einschränkungen
- restricted - sehr streng
- Reference: <https://kubernetes.io/docs/concepts/security/pod-security-standards/>

### Praktisches Beispiel für Version ab 1.2.23 - Problemstellung

```
mkdir -p manifests
cd manifests
mkdir psa
cd psa
nano 01-ns.yml
```

```
## Schritt 1: Namespace anlegen
## vi 01-ns.yml

apiVersion: v1
kind: Namespace
metadata:
  name: test-ns1
  labels:
    # soft version - running but showing complaints
    # pod-security.kubernetes.io/enforce: baseline
    pod-security.kubernetes.io/enforce: restricted
    pod-security.kubernetes.io/audit: restricted
    pod-security.kubernetes.io/warn: restricted
```

```
kubectl apply -f 01-ns.yml
```

```
## Schritt 2: Testen mit nginx - pod
## vi 02-nginx.yml

apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: test-ns1
spec:
  containers:
    - image: nginx
      name: nginx
      ports:
        - containerPort: 80
```

```
## a lot of warnings will come up
## because this image runs as root !! (by default)
kubectl apply -f 02-nginx.yml
```

```
## Schritt 3:
## Anpassen der Sicherheitseinstellung (Phase1) im Container

## vi 02-nginx.yml

apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: test-ns1
spec:
  containers:
    - image: nginx
      name: nginx
      ports:
        - containerPort: 80
      securityContext:
        seccompProfile:
          type: RuntimeDefault
```

```
kubectl delete -f 02-nginx.yml
kubectl apply -f 02-nginx.yml
kubectl -n test-ns1 get pods
```

```
## Schritt 4:
## Weitere Anpassung runAsNotRoot
## vi 02-nginx.yml
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: test-ns<tln>
spec:
  containers:
    - image: nginx
      name: nginx
      ports:
        - containerPort: 80
      securityContext:
        seccompProfile:
          type: RuntimeDefault
      runAsNonRoot: true
```

```
## pod kann erstellt werden, wird aber nicht gestartet
kubectl delete -f 02-nginx.yml
kubectl apply -f 02-nginx.yml
kubectl -n test-ns1 get pods
kubectl -n test-ns1 describe pods nginx
```

```
## Schritt 4:
## Anpassen der Sicherheitseinstellung (Phase1) im Container
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: test-ns1
spec:
  containers:
    - image: nginx
      name: nginx
      ports:
        - containerPort: 80
      securityContext:
        seccompProfile:
          type: RuntimeDefault
        runAsNonRoot: true
        allowPrivilegeEscalation: false
        capabilities:
          drop: ["ALL"]
```

```
kubectl delete -f 02-nginx.yml
kubectl apply -f 02-nginx.yml
kubectl -n test-ns1 get pods
```

#### Praktisches Beispiel für Version ab 1.2.23 -Lösung - Container als NICHT-Root laufen lassen

- Wir müssen ein image, dass auch als NICHT-Root laufen kann
- .. oder selbst eines bauen (:o)) o bei nginx ist das bitnami/nginx

```
## vi 03-nginx-bitnami.yml
apiVersion: v1
kind: Pod
metadata:
  name: bitnami-nginx
  namespace: test-ns1
spec:
  containers:
    - image: bitnami/nginx
      name: bitnami-nginx
      ports:
        - containerPort: 80
      securityContext:
        seccompProfile:
          type: RuntimeDefault
        runAsNonRoot: true
```

```
## und er läuft als nicht root
kubectl apply -f 03_pod-bitnami.yml
kubectl -n test-ns1 get pods
```

## Was muss ich bei der Netzwerk-Sicherheit beachten ?

### Bereich 1: Kubernetes (Cluster)

```
1. Welche Ports sollten wirklich geöffnet sein ?

für Kubernetes

2. Wer muss den von wo den Kube-API-Server zugreifen

- den Traffic einschränken
```

### Bereich 2: Nodes

Alle nicht benötigten fremden Ports sollten geschlossen sein  
Wenn offen, nur über vordefinierte Zugangswege (und auch nur bestimmte Nutzer)

### Pods (Container / Image)

```
## Ingress (NetworkPolicy) - engmaschig stricken
## 1. Wer soll von wo auf welche Pod zugreifen können

## 2. Welche Pod auf welchen anderen Pod (Service)

Egress
## Welche Pods dürfen wohin nach draussen
```

### Einschränkung der Fähigkeiten eines Pods

kein PrivilegeEscalation  
nur notwendige Capabilities  
unter einem nicht-root Benutzer laufen lassen

### Patching

```
## pods -> neuestes images bei security vulnerabilities
## nodes -> auch neues patches (apt upgrade)
## kubernetes cluster -> auf dem neuesten Stand
# -> wie ist der Prozess ClusterUpdate, update der manifeste zu neuen API-Versionen
```

### RBAC

```
## Nutzer (kubectl, systemnutzer -> pods)

## 1. Zugriff von den pods

## 2. Zugriff über helm / kubectl
## Wer darf was ? Was muss der Nutzer können
```

### Compliance

PSP's / PSA  
PodSecurityPolicy was deprecated in Kubernetes v1.21, and removed from Kubernetes in v1.25  
PSA - Pod Security Admission

## Kubernetes Interna / Misc.

### OCI,Container,Images Standards

#### Schritt 1:

```
cd
mkdir bautest
cd bautest
```

#### Schritt 2:

```
## nano docker-compose.yml
version: "3.8"

services:
  myubuntu:
    build: ./myubuntu
    restart: always
```

### Schritt 3:

```
mkdir myubuntu
cd myubuntu
```

```
nano hello.sh
```

```
#!/bin/bash
let i=0

while true
do
  let i=i+1
  echo $i:hello-docker
  sleep 5
done
```

```
## nano Dockerfile
FROM ubuntu:latest
RUN apt-get update; apt-get install -y inetutils-ping
COPY hello.sh .
RUN chmod u+x hello.sh
CMD ["/hello.sh"]
```

### Schritt 4:

```
cd ../
## wichtig, im docker-compose - Ordner seiend
##pwd
##~/bautest
docker-compose up -d
## wird image gebaut und container gestartet

## Bei Veränderung vom Dockerfile, muss man den Parameter --build mitangeben
docker-compose up -d --build
```

### Geolocation Kubernetes Cluster

- <https://learnk8s.io/bite-sized/connecting-multiple-kubernetes-clusters>

## Kubernetes - Überblick

### Installation - Welche Komponenten from scratch

#### Step 1: Server 1 (manuell installiert -> microk8s)

```
## Installation Ubuntu - Server

## cloud-init script
## s.u. BASIS (keine Voraussetzung - nur zum Einrichten des Nutzers 11trainingdo per ssh)

## Server 1 - manuell
## Ubuntu 20.04 LTS - Grundinstallation

## minimal Netzwerk - öffentlichen IP
## nichts besonderes eingerichtet - Standard Digitalocean

## Standard vo Installation microk8s
lo UNKNOWN 127.0.0.1/8 ::1/128
## public ip / interne
eth0 UP 164.92.255.234/20 10.19.0.6/16 fe80::c:66ff:fec4:cbce/64
## private ip
eth1 UP 10.135.0.3/16 fe80::8081:aaff:feaa:780/64

snap install microk8s --classic
## namensauflösung fuer pods
microk8s enable dns
```



```
## Funktioniert microk8s
microk8s status
```

## Steps 2: Server 2+3 (automatische Installation -> microk8s )

```
## Was macht das ?
## 1. Basisnutzer (11trainingdo) - keine Voraussetzung für microk8s
## 2. Installation von microk8s
##.>>>>>>> microk8s installiert <<<<<<<<
## - snap install --classic microk8s
## >>>>>>> Zuordnung zur Gruppe microk8s - notwendig für bestimmte plugins (z.B. helm)
## usermod -a -G microk8s root
## >>>>>>> Setzen des .kube - Verzeichnisses auf den Nutzer microk8s -> nicht zwingend erforderlich
## chown -R -R microk8s ~/.kube
## >>>>>>> REQUIRED .. DNS aktivieren, wichtig für Namensauflösungen innerhalb der PODS
## >>>>>>> sonst funktioniert das nicht !!!
## microk8s enable dns
## >>>>>>> kubectl alias gesetzt, damit man nicht immer microk8s kubectl eingeben muss
## - echo "alias kubectl='microk8s kubectl'" >> /root/.bashrc

## cloud-init script
## s.u. MITMICROK8S (keine Voraussetzung - nur zum Einrichten des Nutzers 11trainingdo per ssh)
##cloud-config
users:
  - name: 11trainingdo
    shell: /bin/bash

runcmd:
  - sed -i "s/PasswordAuthentication no/PasswordAuthentication yes/g" /etc/ssh/sshd_config
  - echo " " >> /etc/ssh/sshd_config
  - echo "AllowUsers 11trainingdo" >> /etc/ssh/sshd_config
  - echo "AllowUsers root" >> /etc/ssh/sshd_config
  - systemctl reload sshd
  - sed -i '/11trainingdo/c
11trainingdo:$6$HeLUjW3a$4xSfDFQjKWfAoGkZF3LFAxM4hgl3d6ATbr2kEu9zMOfwLxkYMO.AJF526mZONwdmsm9sg0tCBKl.SYbhS52u70:17476:0:99999:7:::
/etc/shadow
  - echo "11trainingdo ALL=(ALL) ALL" > /etc/sudoers.d/11trainingdo
  - chmod 0440 /etc/sudoers.d/11trainingdo

  - echo "Installing microk8s"
  - snap install --classic microk8s
  - usermod -a -G microk8s root
  - chown -f -R microk8s ~/.kube
  - microk8s enable dns
  - echo "alias kubectl='microk8s kubectl'" >> /root/.bashrc

## Prüfen ob microk8s - wird automatisch nach Installation gestartet
## kann eine Weile dauern
microk8s status
```

## Step 3: Client - Maschine (wir sollten nicht auf control-plane oder cluster - node arbeiten

```
Weiteren Server hochgezogen.
Vanilla + BASIS

## Installation Ubuntu - Server

## cloud-init script
## s.u. BASIS (keine Voraussetzung - nur zum Einrichten des Nutzers 11trainingdo per ssh)

## Server 1 - manuell
## Ubuntu 20.04 LTS - Grundinstallation

## minimal Netzwerk - öffentlichen IP
## nichts besonderes eingerichtet - Standard Digitalocean

## Standard vo Installation microk8s
lo                UNKNOWN          127.0.0.1/8 ::1/128
## public ip / interne
eth0              UP                164.92.255.232/20 10.19.0.6/16 fe80::c:66ff:fec4:cbce/64
## private ip
eth1              UP                10.135.0.5/16 fe80::8081:aaff:feaa:780/64

#### Installation von kubectl aus dem snap
## NICHT .. keine microk8s - keine control-plane / worker-node
## NUR Client zum Arbeiten
```

```

snap install kubect1 --classic

#### .kube/config
## Damit ein Zugriff auf die kube-server-api möglich
## d.h. REST-API Interface, um das Cluster verwalten.
## Hier haben uns für den ersten Control-Node entschieden
## Alternativ wäre round-robin per dns möglich

## Mini-Schritt 1:
## Auf dem Server 1: kubeconfig ausspielen
microk8s config > /root/kube-config
## auf das Zielsystem gebracht (client 1)
scp /root/kubeconfig 11trainingdo@10.135.0.5:/home/11trainingdo

## Mini-Schritt 2:
## Auf dem Client 1 (diese Maschine) kubeconfig an die richtige Stelle bringen
## Standardmäßig der Client nach eine Konfigurationsdatei sucht in ~/.kube/config
sudo su -
cd
mkdir .kube
cd .kube
mv /home/11trainingdo/kube-config config

## Verbindungstest gemacht
## Damit feststellen ob das funktioniert.
kubect1 cluster-info

```

#### Schritt 4: Auf allen Servern IP's hinterlegen und richtigen Hostnamen überprüfen

```

## Auf jedem Server
hostnamectl
## evtl. hostname setzen
## z.B. - auf jedem Server eindeutig
hostnamectl set-hostname n1.training.local

## Gleiche hosts auf allen server einrichten.
## Wichtig, um Traffic zu minimieren verwenden, die interne (private) IP

/etc/hosts
10.135.0.3 n1.training.local n1
10.135.0.4 n2.training.local n2
10.135.0.5 n3.training.local n3

```

#### Schritt 5: Cluster aufbauen

```

## Mini-Schritt 1:
## Server 1: connection - string (token)
microk8s add-node
## Zeigt Liste und wir nehmen den Eintrag mit der lokalen / öffentlichen ip
## Dieser Token kann nur 1x verwendet werden und wir auf dem ANDEREN node ausgeführt
## microk8s join 10.135.0.3:25000/e9cdaa11b5d6d24461c8643cdf107837/bcad1949221a

## Mini-Schritt 2:
## Dauert eine Weile, bis das durch ist.
## Server 2: Den Node hinzufügen durch den JOIN - Befehl
microk8s join 10.135.0.3:25000/e9cdaa11b5d6d24461c8643cdf107837/bcad1949221a

## Mini-Schritt 3:
## Server 1: token besorgen für node 3
microk8s add-node

## Mini-Schritt 4:
## Server 3: Den Node hinzufügen durch den JOIN-Befehl
microk8s join 10.135.0.3:25000/09c96e57ec12af45b2752fb45450530c/bcad1949221a

## Mini-Schritt 5: Überprüfen ob HA-Cluster läuft
Server 1: (es kann auf jedem der 3 Server überprüft werden, auf einem reicht
microk8s status | grep high-availability
high-availability: yes

```

#### Ergänzend nicht notwendige Scripte

```

## cloud-init script
## s.u. BASIS (keine Voraussetzung - nur zum Einrichten des Nutzers 11trainingdo per ssh)

## Digitalocean - unter user_data reingepastet beim Einrichten

```

```
##cloud-config
users:
  - name: 11trainingdo
    shell: /bin/bash

runcmd:
  - sed -i "s/PasswordAuthentication no/PasswordAuthentication yes/g" /etc/ssh/sshd_config
  - echo " " >> /etc/ssh/sshd_config
  - echo "AllowUsers 11trainingdo" >> /etc/ssh/sshd_config
  - echo "AllowUsers root" >> /etc/ssh/sshd_config
  - systemctl reload sshd
  - sed -i '/11trainingdo/c
11trainingdo:$6$HeLUJW3a$4xSfDFQjKWfAoGkZF3LFAxM4hgl3d6ATbr2kEu9zMOfwLxkYMO.AJF526mZONwdmsm9sg0tCBKl.SYbhS52u70:17476:0:99999:7:::
/etc/shadow
  - echo "11trainingdo ALL=(ALL) ALL" > /etc/sudoers.d/11trainingdo
  - chmod 0440 /etc/sudoers.d/11trainingdo
```

## Kubernetes - microk8s (Installation und Management)

### kubectl unter windows - Remote-Verbindung zu KubereneTs (microk8s) einrichten

#### Walkthrough (Installation)

```
## Step 1
chocolatey installiert.
(powershell als Administrator ausführen)

## https://docs.chocolatey.org/en-us/choco/setup
Set-ExecutionPolicy Bypass -Scope Process -Force; [System.Net.ServicePointManager]::SecurityProtocol =
[System.Net.ServicePointManager]::SecurityProtocol -bor 3072; iex ((New-Object
System.Net.WebClient).DownloadString('https://community.chocolatey.org/install.ps1'))

## Step 2
choco install kubernetes-cli

## Step 3
testen:
kubectl version --client

## Step 4:
## powershell als normaler benutzer öffnen
```

#### Walkthrough (autocompletion)

```
in powershell (normaler Benutzer)
kubectl completion powershell | Out-String | Invoke-Expression
```

#### kubectl - config - Struktur vorbereiten

```
## in powershell im heimatordner des Benutzers .kube - ordnern anlegen
## C:\Users\<dein-name>\
mkdir .kube
cd .kube
```

#### IP von Cluster-Node bekommen

```
## auf virtualbox - maschine per ssh einloggen
## öffentliche ip herausfinden - z.B. enp0s8 bei HostOnly - Adapter
ip -br a
```

#### config für kubectl aus Cluster-Node auslesen (microk8s)

```
## auf virtualbox - maschine per ssh einloggen / zum root wechseln
## abfragen
microk8s config

## Alle Zeilen ins clipboard kopieren
## und mit notepad++ in die Datei \Users\<dein-name>\.kube\config
## schreiben

## Wichtig: Zeile cluster -> clusters / server
## Hier ip von letztem Schritt eintragen:
## z.B.
Server: https://192.168.56.106/.....
```

## Testen

```
## in powershell
## kann ich eine Verbindung zum Cluster aufbauen ?
kubectl cluster-info
```

- <https://kubernetes.io/docs/tasks/tools/install-kubectl-windows/>

## Arbeiten mit der Registry

### Installation Kubernetes Dashboard

#### Reference:

- <https://blog.tippyybits.com/installing-kubernetes-in-virtualbox-3d49f666b4d6>

## Kubernetes - RBAC

### Nutzer einrichten - kubernetes bis 1.24

#### Enable RBAC in microk8s

```
## This is important, if not enable every user on the system is allowed to do everything
microk8s enable rbac
```

#### Schritt 1: Nutzer-Account auf Server anlegen / in Client

```
cd
mkdir -p manifests/rbac
cd manifests/rbac
```

##### Mini-Schritt 1: Definition für Nutzer

```
## vi service-account.yml
apiVersion: v1
kind: ServiceAccount
metadata:
  name: training
  namespace: default
```

```
kubectl apply -f service-account.yml
```

##### Mini-Schritt 2: ClusterRolle festlegen - Dies gilt für alle namespaces, muss aber noch zugewiesen werden

```
### Bevor sie zugewiesen ist, funktioniert sie nicht - da sie keinem Nutzer zugewiesen ist
```

```
## vi pods-clusterrole.yml
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: pods-clusterrole
rules:
- apiGroups: [""] # "" indicates the core API group
  resources: ["pods"]
  verbs: ["get", "watch", "list"]
```

```
kubectl apply -f pods-clusterrole.yml
```

##### Mini-Schritt 3: Die ClusterRolle den entsprechenden Nutzern über RoleBinding zu ordnen

```
## vi rb-training-ns-default-pods.yml
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: rolebinding-ns-default-pods
  namespace: default
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: pods-clusterrole
subjects:
- kind: ServiceAccount
  name: training
  namespace: default
```

```
kubectl apply -f rb-training-ns-default-pods.yml
```

#### Mini-Schritt 4: Testen (klappt der Zugang)

```
kubectl auth can-i get pods -n default --as system:serviceaccount:default:training
```

### Schritt 2: Context anlegen / Credentials auslesen und in kubeconfig hinterlegen (bis Version 1.25.)

#### Mini-Schritt 1: kubeconfig setzen

```
kubectl config set-context training-ctx --cluster microk8s-cluster --user training

## extract name of the token from here

TOKEN=$(kubectl get secret trainingtoken -o jsonpath='{.data.token}' | base64 --decode)
echo $TOKEN
kubectl config set-credentials training --token=$TOKEN
kubectl config use-context training-ctx

## Hier reichen die Rechte nicht aus
kubectl get deploy
## Error from server (Forbidden): pods is forbidden: User "system:serviceaccount:kube-system:training" cannot list # resource
"pods" in API group "" in the namespace "default"
```

#### Mini-Schritt 2:

```
kubectl config use-context training-ctx
kubectl get pods
```

#### Refs:

- <https://docs.oracle.com/en-us/iaas/Content/ContEng/Tasks/contengaddingserviceaccttoken.htm>
- <https://microk8s.io/docs/multi-user>
- <https://faun.pub/kubernetes-rbac-use-one-role-in-multiple-namespaces-d1d08bb08286>

#### Ref: Create Service Account Token

- <https://kubernetes.io/docs/reference/access-authz/service-accounts-admin/#create-token>

## kubectl

### Tipps&Tricks zu Deployment - Rollout

#### Warum

```
Rückgängig machen von deploys, Deploys neu unstossen.
(Das sind die wichtigsten Fähigkeiten)
```

#### Beispiele

```
## Deployment nochmal durchführen
## z.B. nach kubectl uncordon n12.training.local
kubectl rollout restart deploy nginx-deployment

## Rollout rückgängig machen
kubectl rollout undo deploy nginx-deployment
```

## Kubernetes - Monitoring (microk8s und vanilla)

### metrics-server aktivieren (microk8s und vanilla)

#### Warum ? Was macht er ?

```
Der Metrics-Server sammelt Informationen von den einzelnen Nodes und Pods
Er bietet mit

kubectl top pods
kubectl top nodes

ein einfaches Interface, um einen ersten Eindruck über die Auslastung zu bekommen.
```

#### Walkthrough

```
## Auf einem der Nodes im Cluster (HA-Cluster)
microk8s enable metrics-server

## Es dauert jetzt einen Moment bis dieser aktiv ist auch nach der Installation
## Auf dem Client
```

```
kubectl top nodes
kubectl top pods
```

## Kubernetes

- <https://kubernetes-sigs.github.io/metrics-server/>
- kubectl apply -f <https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml>

## Kubernetes - Backups

## Kubernetes - Tipps & Tricks

### Assigning Pods to Nodes

#### Walkthrough

```
## leave n3 as is
kubectl label nodes n7 rechenzentrum=rz1
kubectl label nodes n17 rechenzentrum=rz2
kubectl label nodes n27 rechenzentrum=rz2

kubectl get nodes --show-labels

## nginx-deployment
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 9 # tells deployment to run 2 pods matching the template
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
      nodeSelector:
        rechenzentrum: rz2

## Let's rewrite that to deployment
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    env: test
spec:
  containers:
    - name: nginx
      image: nginx
      imagePullPolicy: IfNotPresent
  nodeSelector:
    rechenzentrum=rz2
```

#### Ref:

- <https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/>

## Kubernetes - Documentation

### LDAP-Anbindung

- <https://github.com/apprenda-kismatic/kubernetes-ldap>

### Helpful to learn - Kubernetes

- <https://kubernetes.io/docs/tasks/>

### Environment to learn

- <https://killercoda.com/killer-shell-cks>

## Environment to learn II

- <https://killercoda.com/>

## Youtube Channel

- <https://www.youtube.com/watch?v=01gcYSck1c4>

## Kubernetes - Shared Volumes

### Shared Volumes with nfs

#### Create new server and install nfs-server

```
## on Ubuntu 20.04LTS
apt install nfs-kernel-server
systemctl status nfs-server

vi /etc/exports
## adjust ip's of kubernetes master and nodes
## kmaster
/var/nfs/ 192.168.56.101(rw,sync,no_root_squash,no_subtree_check)
## knode1
/var/nfs/ 192.168.56.103(rw,sync,no_root_squash,no_subtree_check)
## knode 2
/var/nfs/ 192.168.56.105(rw,sync,no_root_squash,no_subtree_check)

exportfs -av
```

#### On all nodes (needed for production)

```
##
apt install nfs-common
```

#### On all nodes (only for testing)

```
#### Please do this on all servers (if you have access by ssh)
### find out, if connection to nfs works !

## for testing
mkdir /mnt/nfs
## 10.135.0.18 is our nfs-server
mount -t nfs 10.135.0.18:/var/nfs /mnt/nfs
ls -la /mnt/nfs
umount /mnt/nfs
```

#### Persistent Storage-Step 1: Setup PersistentVolume in cluster

```
cd
cd manifests
mkdir -p nfs
cd nfs
nano 01-pv.yml
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  # any PV name
  name: pv-nfs-tln<nr>
  labels:
    volume: nfs-data-volume-tln<nr>
spec:
  capacity:
    # storage size
    storage: 1Gi
  accessModes:
    # ReadWriteMany(RW from multi nodes), ReadWriteOnce(RW from a node), ReadOnlyMany(R from multi nodes)
    - ReadWriteMany
  persistentVolumeReclaimPolicy:
    # retain even if pods terminate
    Retain
  nfs:
    # NFS server's definition
    path: /var/nfs/tln<nr>/nginx
    server: 10.135.0.18
    readOnly: false
  storageClassName: ""
```

```
kubectl apply -f 01-pv.yml
kubectl get pv
```

## Persistent Storage-Step 2: Create Persistent Volume Claim

```
nano 02-pvc.yml
```

```
## vi 02-pvc.yml
## now we want to claim space
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pv-nfs-claim-tln<nr>
spec:
  storageClassName: ""
  volumeName: pv-nfs-tln<nr>
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 1Gi
```

```
kubectl apply -f 02-pvc.yml
kubectl get pvc
```

## Persistent Storage-Step 3: Deployment

```
## deployment including mount
## vi 03-deploy.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 4 # tells deployment to run 4 pods matching the template
  template:
    metadata:
      labels:
        app: nginx
    spec:

      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80

      volumeMounts:
        - name: nfsvol
          mountPath: "/usr/share/nginx/html"

      volumes:
        - name: nfsvol
          persistentVolumeClaim:
            claimName: pv-nfs-claim-tln<tln>
```

```
kubectl apply -f 03-deploy.yml
```

## Persistent Storage Step 4: service

```
## now testing it with a service
## cat 04-service.yml
apiVersion: v1
kind: Service
metadata:
  name: service-nginx
  labels:
    run: svc-my-nginx
spec:
  type: NodePort
  ports:
```



```
- port: 80
  protocol: TCP
  selector:
    app: nginx
```

```
kubectl apply -f 04-service.yml
```

### Persistent Storage Step 5: write data and test

```
## connect to the container and add index.html - data
kubectl exec -it deploy/nginx-deployment -- bash
## in container
echo "hello dear friend" > /usr/share/nginx/html/index.html
exit

## now try to connect
kubectl get svc

## connect with ip and port
kubectl run -it --rm curly --image=curlimages/curl -- /bin/sh
## curl http://<cluster-ip>
## exit

## now destroy deployment
kubectl delete -f 03-deploy.yml

## Try again - no connection
kubectl run -it --rm curly --image=curlimages/curl -- /bin/sh
## curl http://<cluster-ip>
## exit
```

### Persistent Storage Step 6: retest after redeployment

```
## now start deployment again
kubectl apply -f 03-deploy.yml

## and try connection again
kubectl run -it --rm curly --image=curlimages/curl -- /bin/sh
## curl http://<cluster-ip>
## exit
```

## Kubernetes - Hardening

### Kubernetes Tipps Hardening

#### PSA (Pod Security Admission)

Policies defined by namespace.  
e.g. not allowed to run container as root.

Will complain/deny when creating such a pod with that container type

### Möglichkeiten in Pods und Containern

```
## für die Pods
kubectl explain pod.spec.securityContext
kubectl explain pod.spec.containers.securityContext
```

### Example (seccomp / security context)

A. seccomp - profile  
<https://github.com/docker/docker/blob/master/profiles/seccomp/default.json>

```
apiVersion: v1
kind: Pod
metadata:
  name: audit-pod
  labels:
    app: audit-pod
spec:
  securityContext:
    seccompProfile:
      type: Localhost
      localhostProfile: profiles/audit.json
```

```
containers:

- name: test-container
  image: hashicorp/http-echo:0.2.3
  args:
  - "-text=just made some syscalls!"
  securityContext:
    allowPrivilegeEscalation: false
```

### SecurityContext (auf Pod Ebene)

```
kubectl explain pod.spec.containers.securityContext
```

### NetworkPolicy

```
## Firewall Kubernetes
```